

RBR

Account(cid , bal)

Set bal of cid 101 , cid 102 to 5000

In SQL:

begin
update Account set
bal = 5000 where cid = 101

update Account set
bal = 5000 where cid = 102

end.

low level operation:

begin transaction T_1

write (bal: 5000 where Cd = 101) → $\omega(\text{Sal})$
write (bal: 5000 where Cd = 102) → $\omega(\text{Sal})$

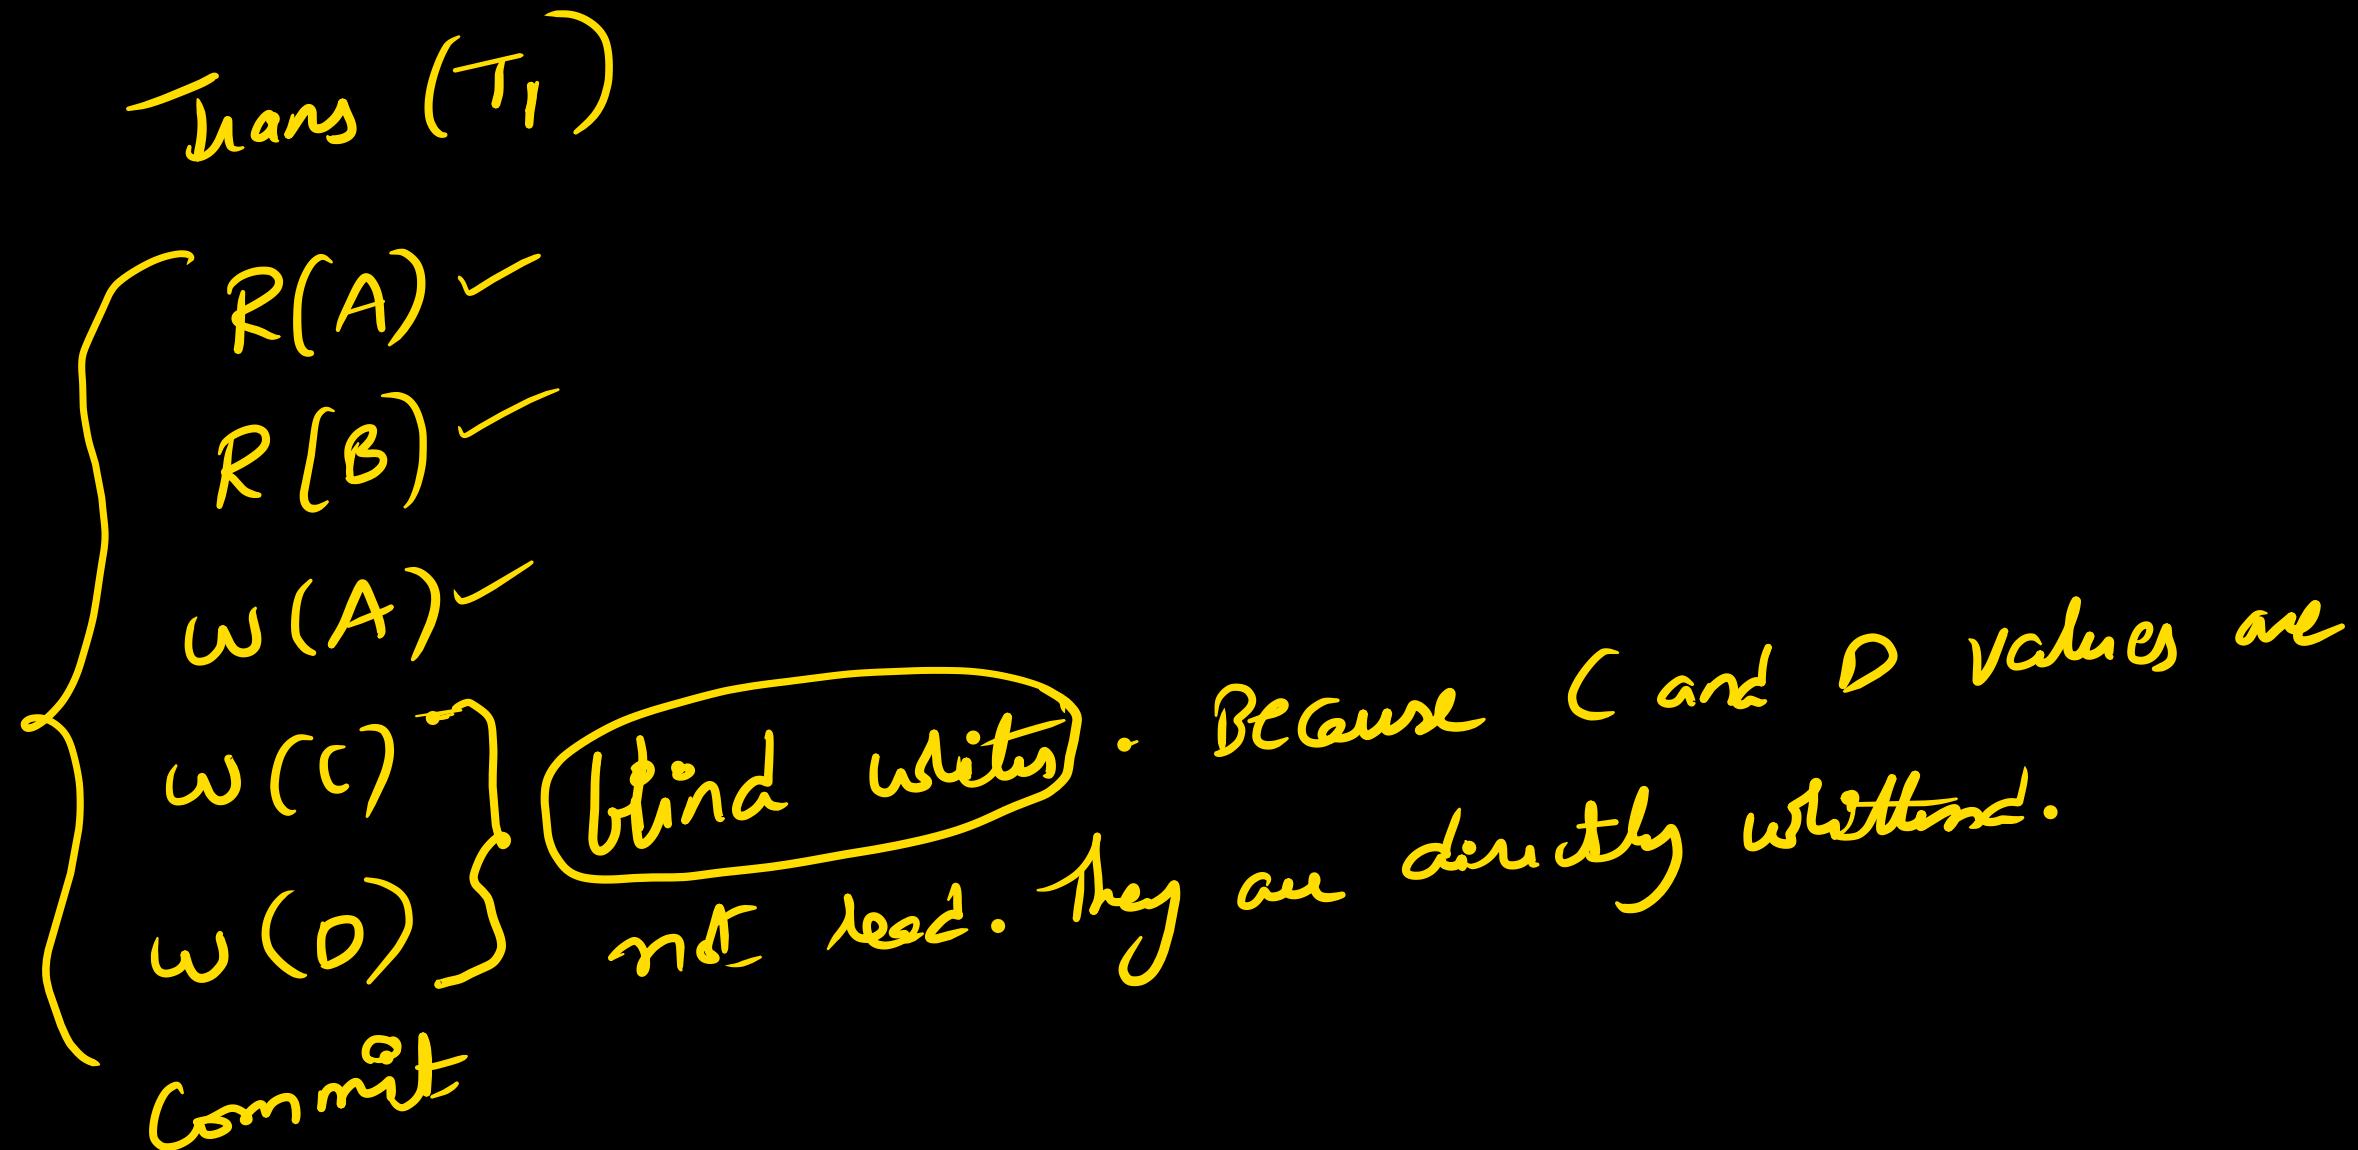
commit

end transaction

✓ Read (sal) → $\omega \& r(\text{Sal})$

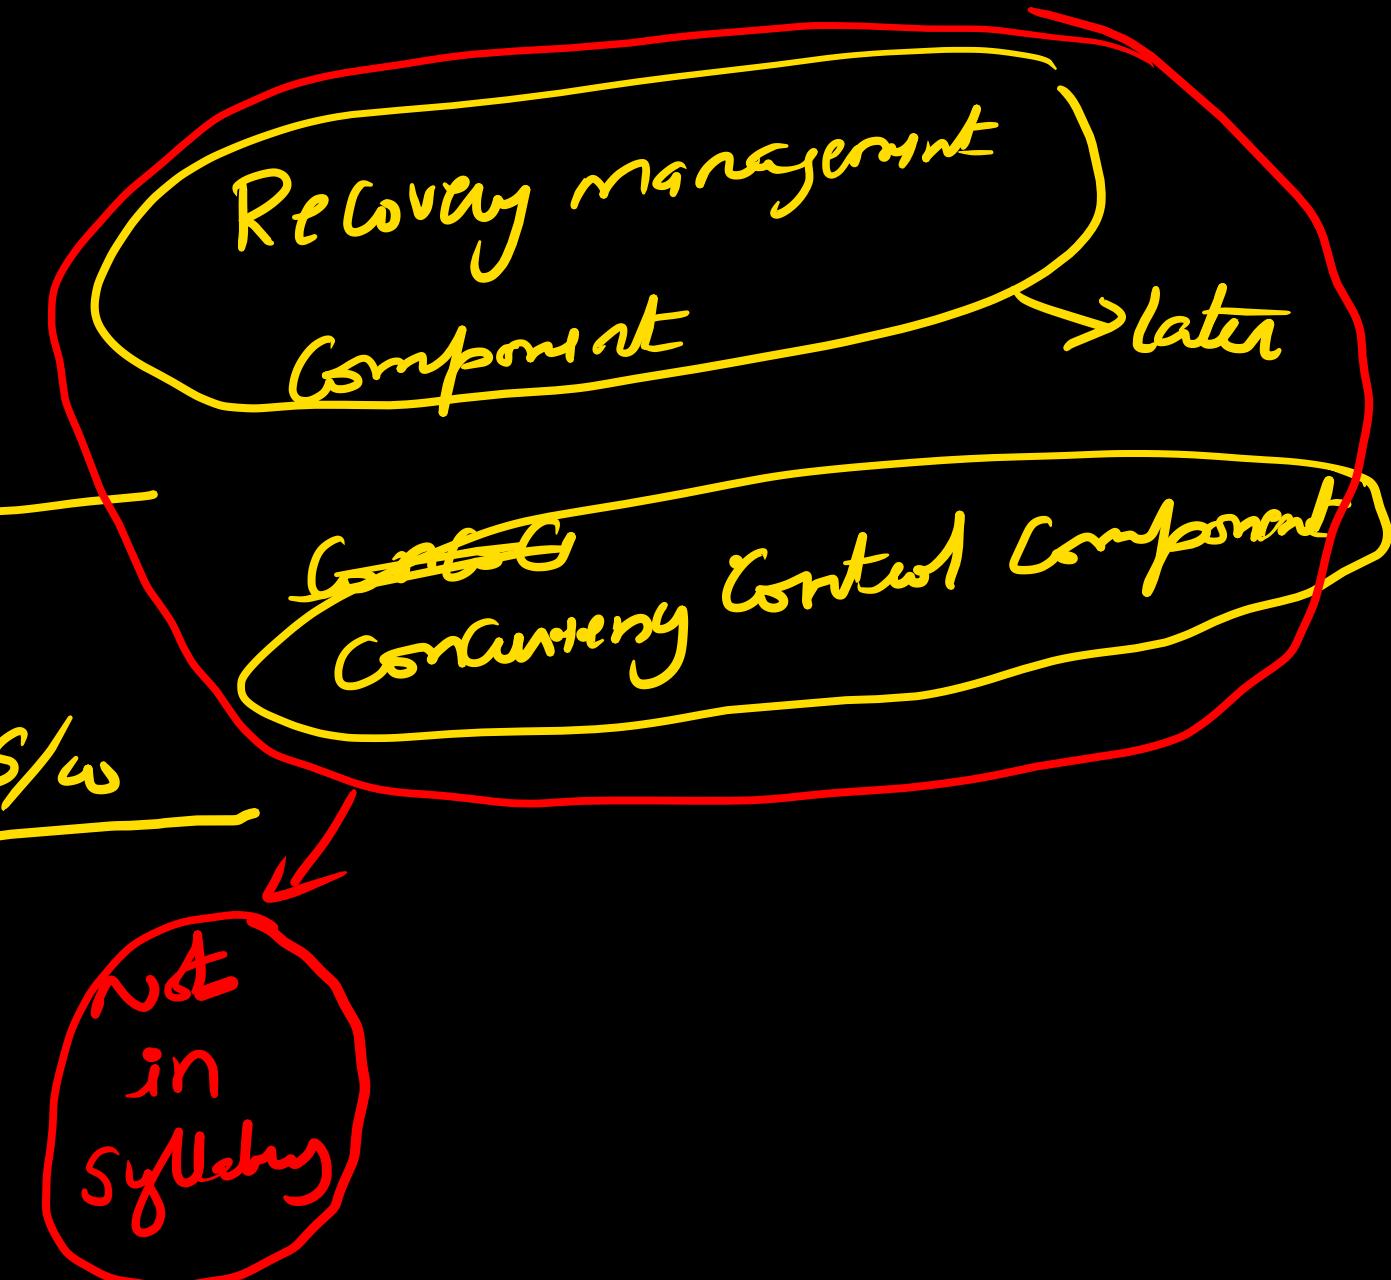
~~Sal = S + 100~~
✓ write (sal) → $\omega(\text{Sal})$

Ex of a transaction:



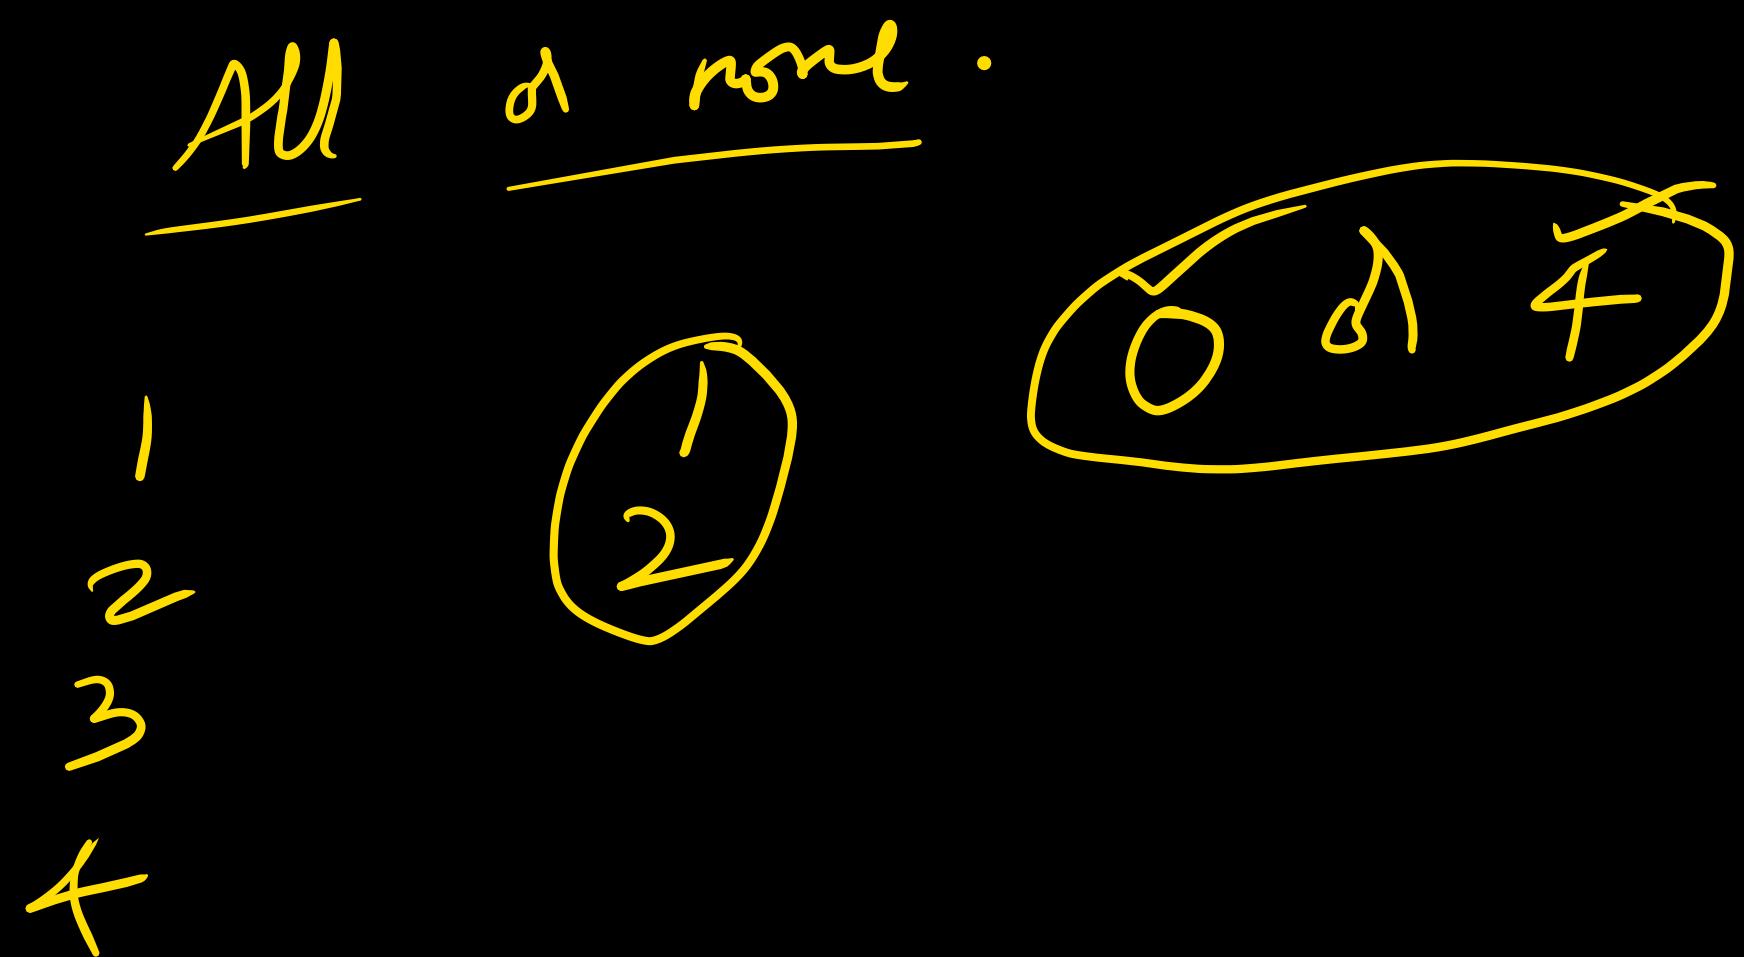
ACID properties: To preserve integrity (correctness) of transactions

- ✓ A - Atomicity
 - ✓ D - Durability
 - ✓ I - Isolation
 - ✓ C - Consistency
- .. } maintained by recovery management component & DBMS S/w
- .. } maintained by concurrency control component & DBMS S/w
- .. } maintained by user DBA / DB dev / DB user.

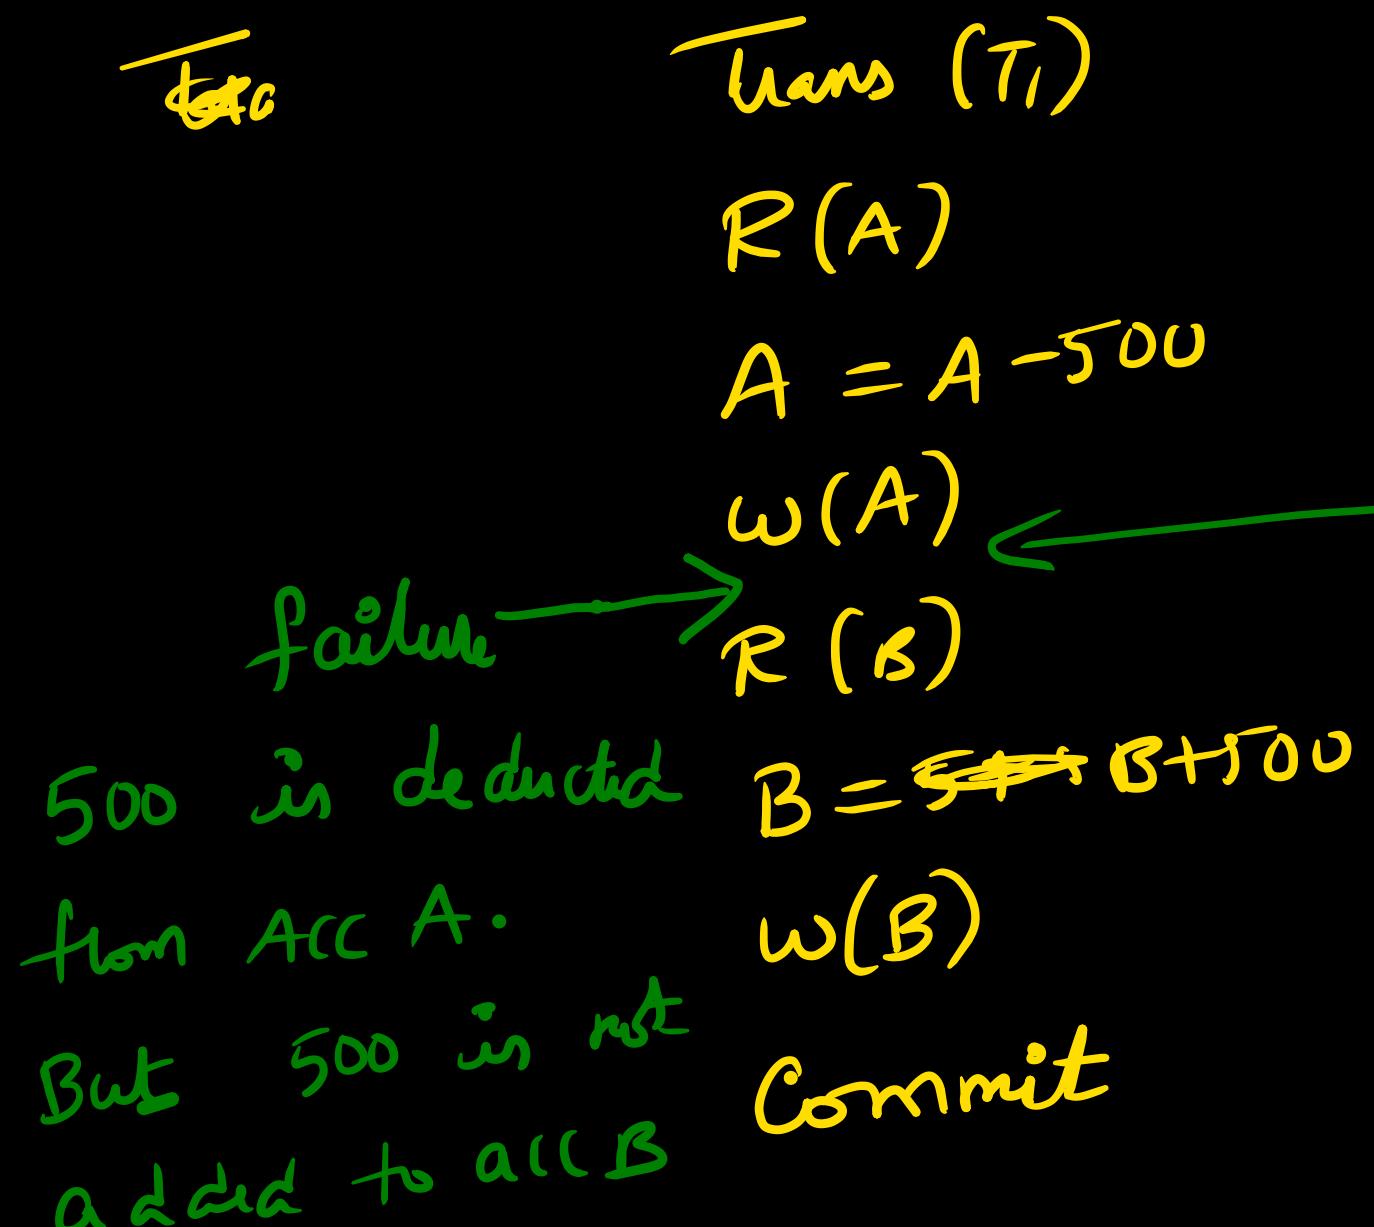


Atomicity: Executes all operations of transaction including commit

(Δ) executes none of the operations of transaction



Ex: There is a transaction to transfer 500 from acc A to acc B.



failure leads to incorrect data.
To preserve integrity we will used
atomicity and do a roll back.

Recovery manager is responsible for successful roll back of a failed transaction.

Recovery management Component:
It rolls back transaction if transaction failed anywhere before commit.

Roll back (or abort): undo modifications of DB file which are done by failed transaction

Q) How do recovery manager perform rollback?

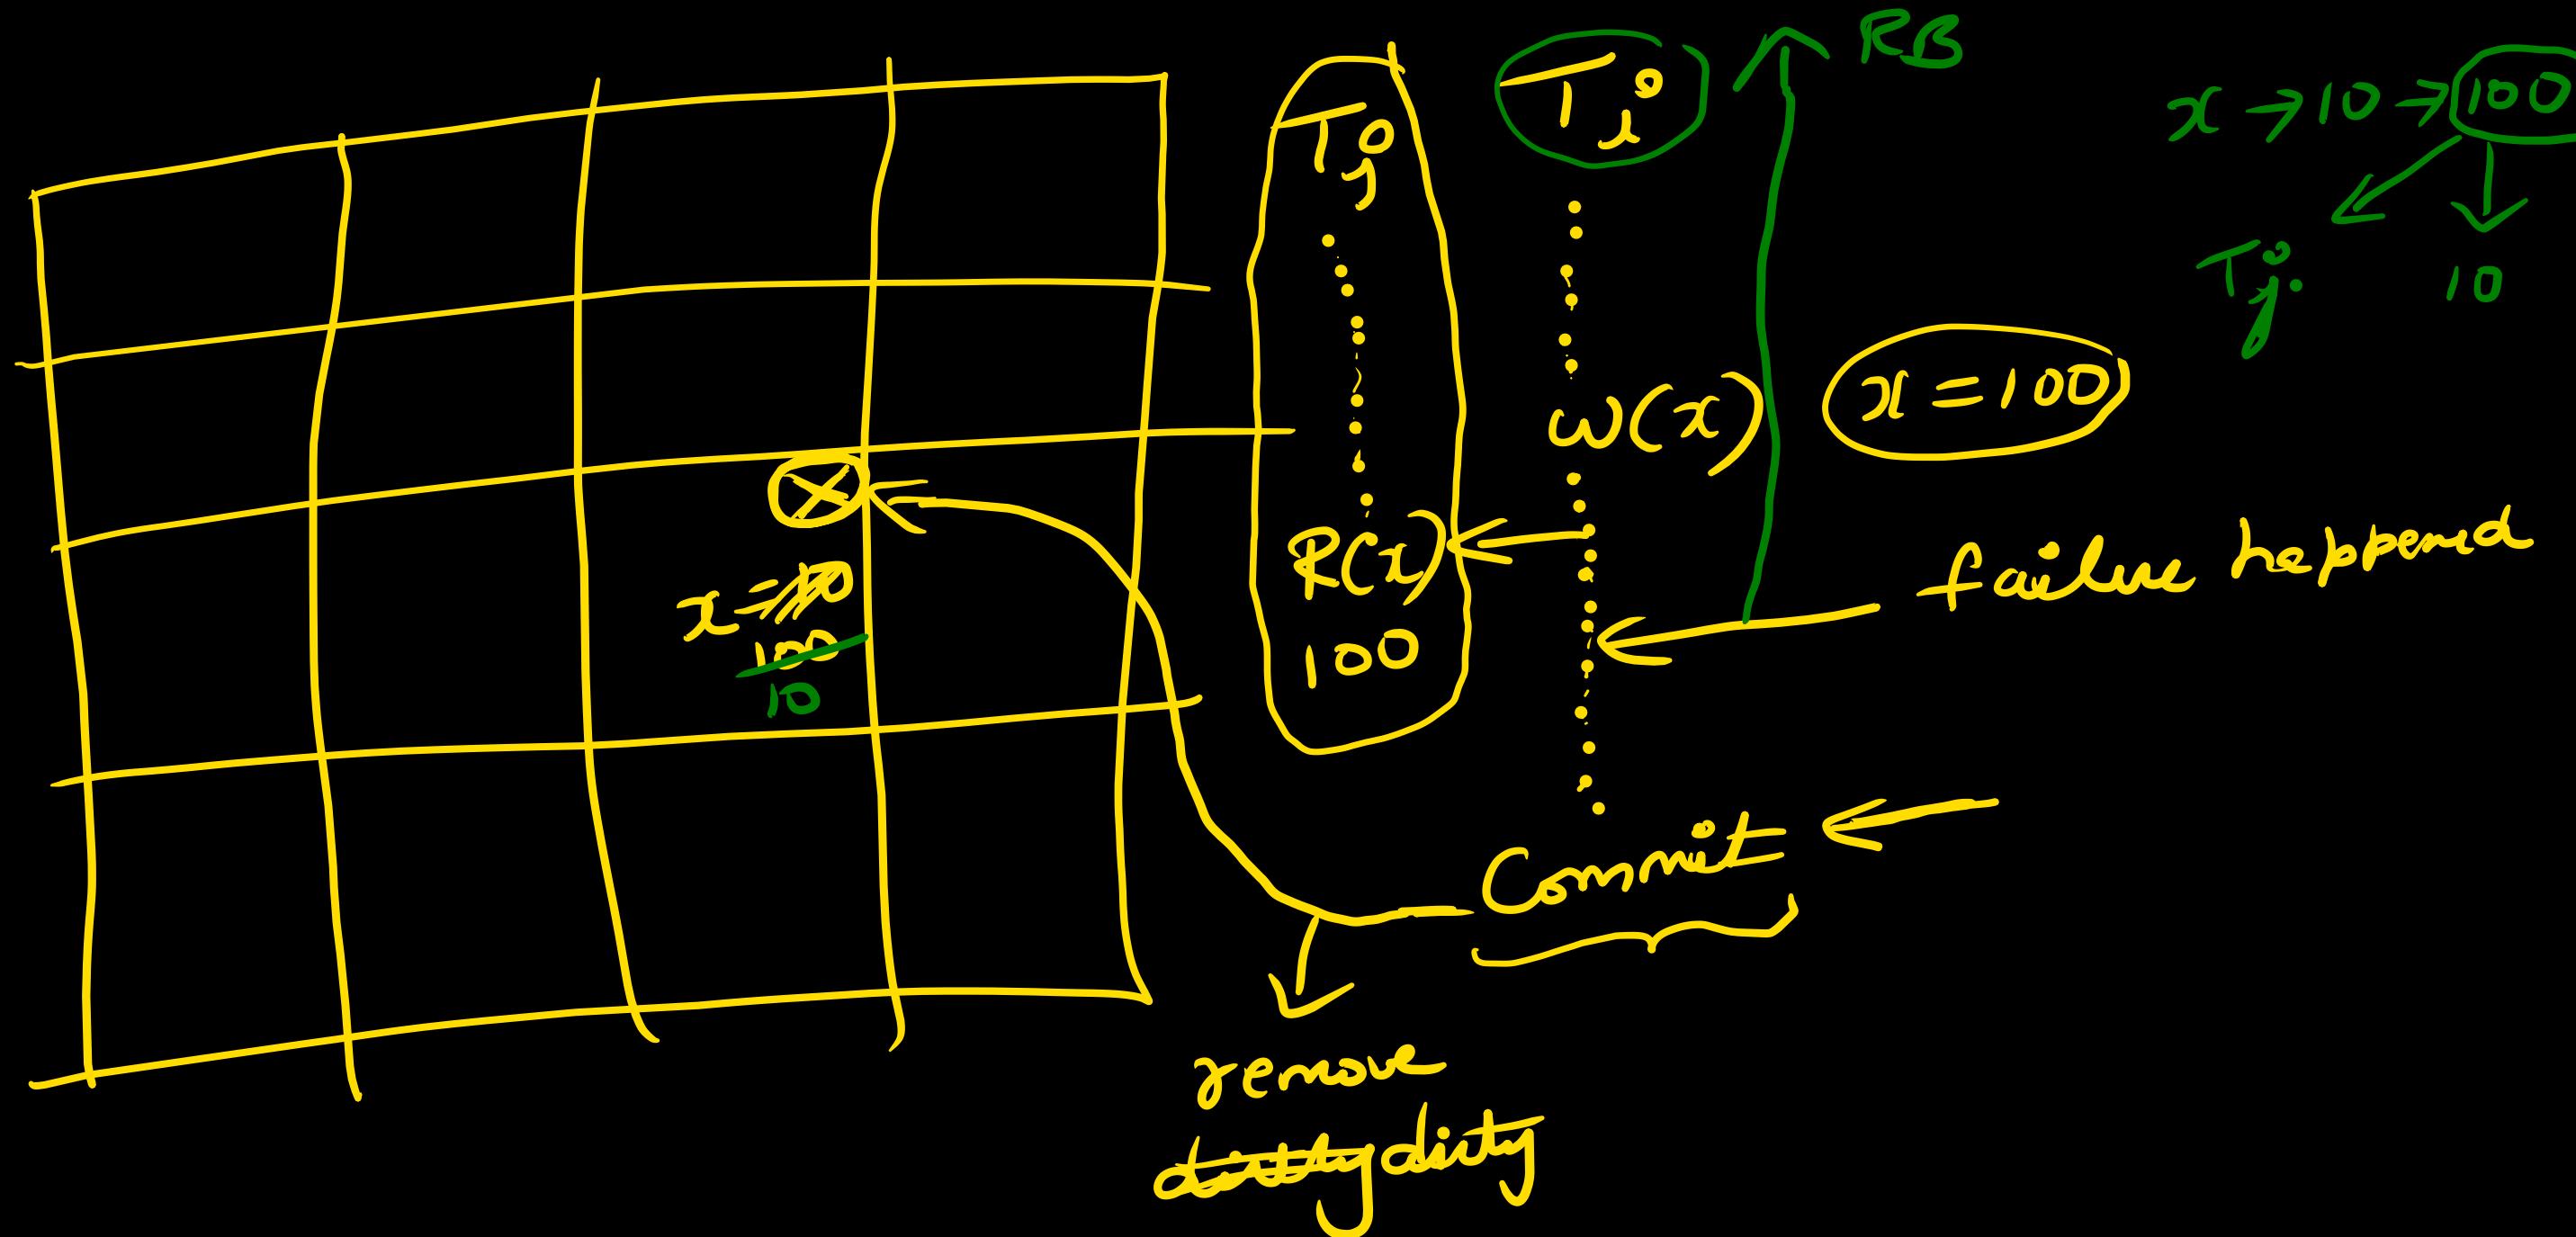
A) Transaction log

Transaction log:

Transaction log

1.	T ₁ begin	every action of any transaction is logged into transaction log
2.	T ₁ , R, A, 100	
3	T ₂ begin	
4	T ₂ , W, B 1000	

Transaction log is a file maintained by recovery manager in the Disk to record all actions of transactions. For roll back, Recovery management component will see transaction log and do it.



A block that is updated by a uncommitted trans is a dirty block.

Why is it dirty? Because there is a chance of rollback and some other transaction might read the uncommitted data. ~~the~~

T_1 (we should not let any transaction
 $R(w)$ access dirty data)
 $CR(w)$
 $R(x)$
 $w(x)$

Commit ← once commit is reached
there is no rollback

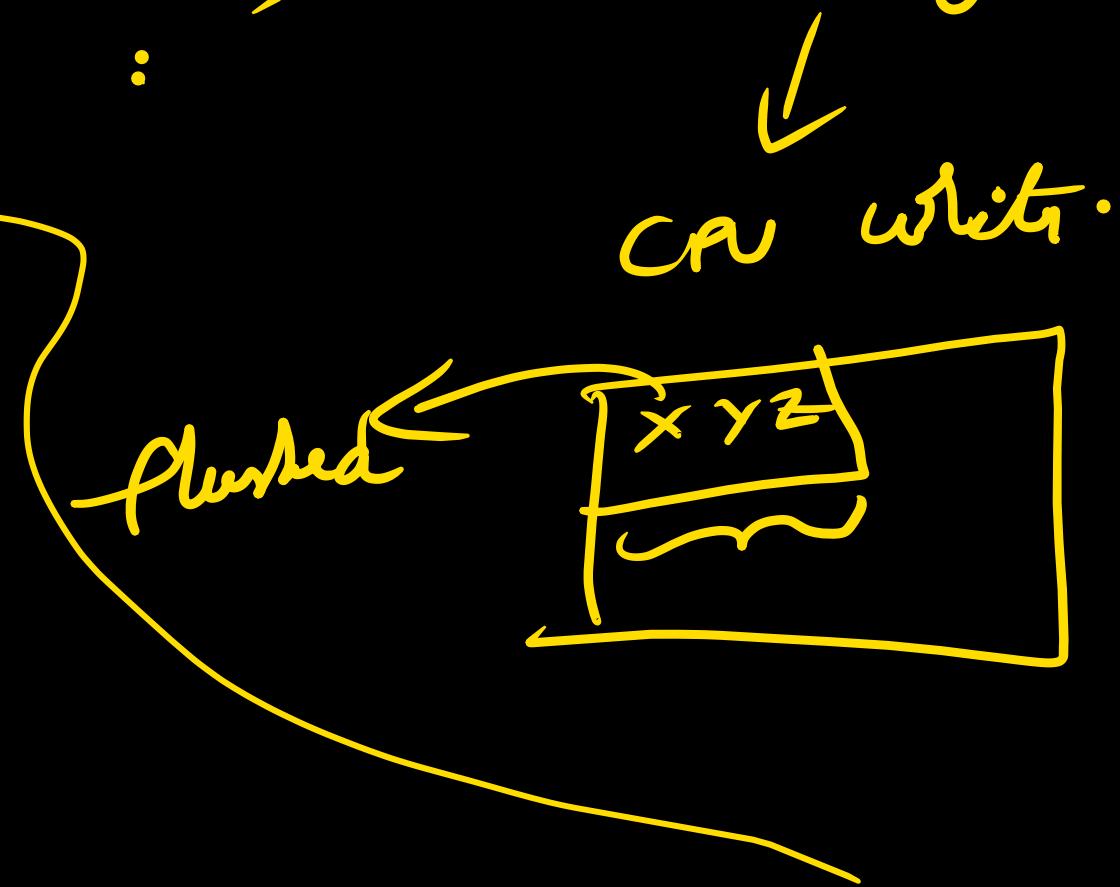
✓ Redo operations:

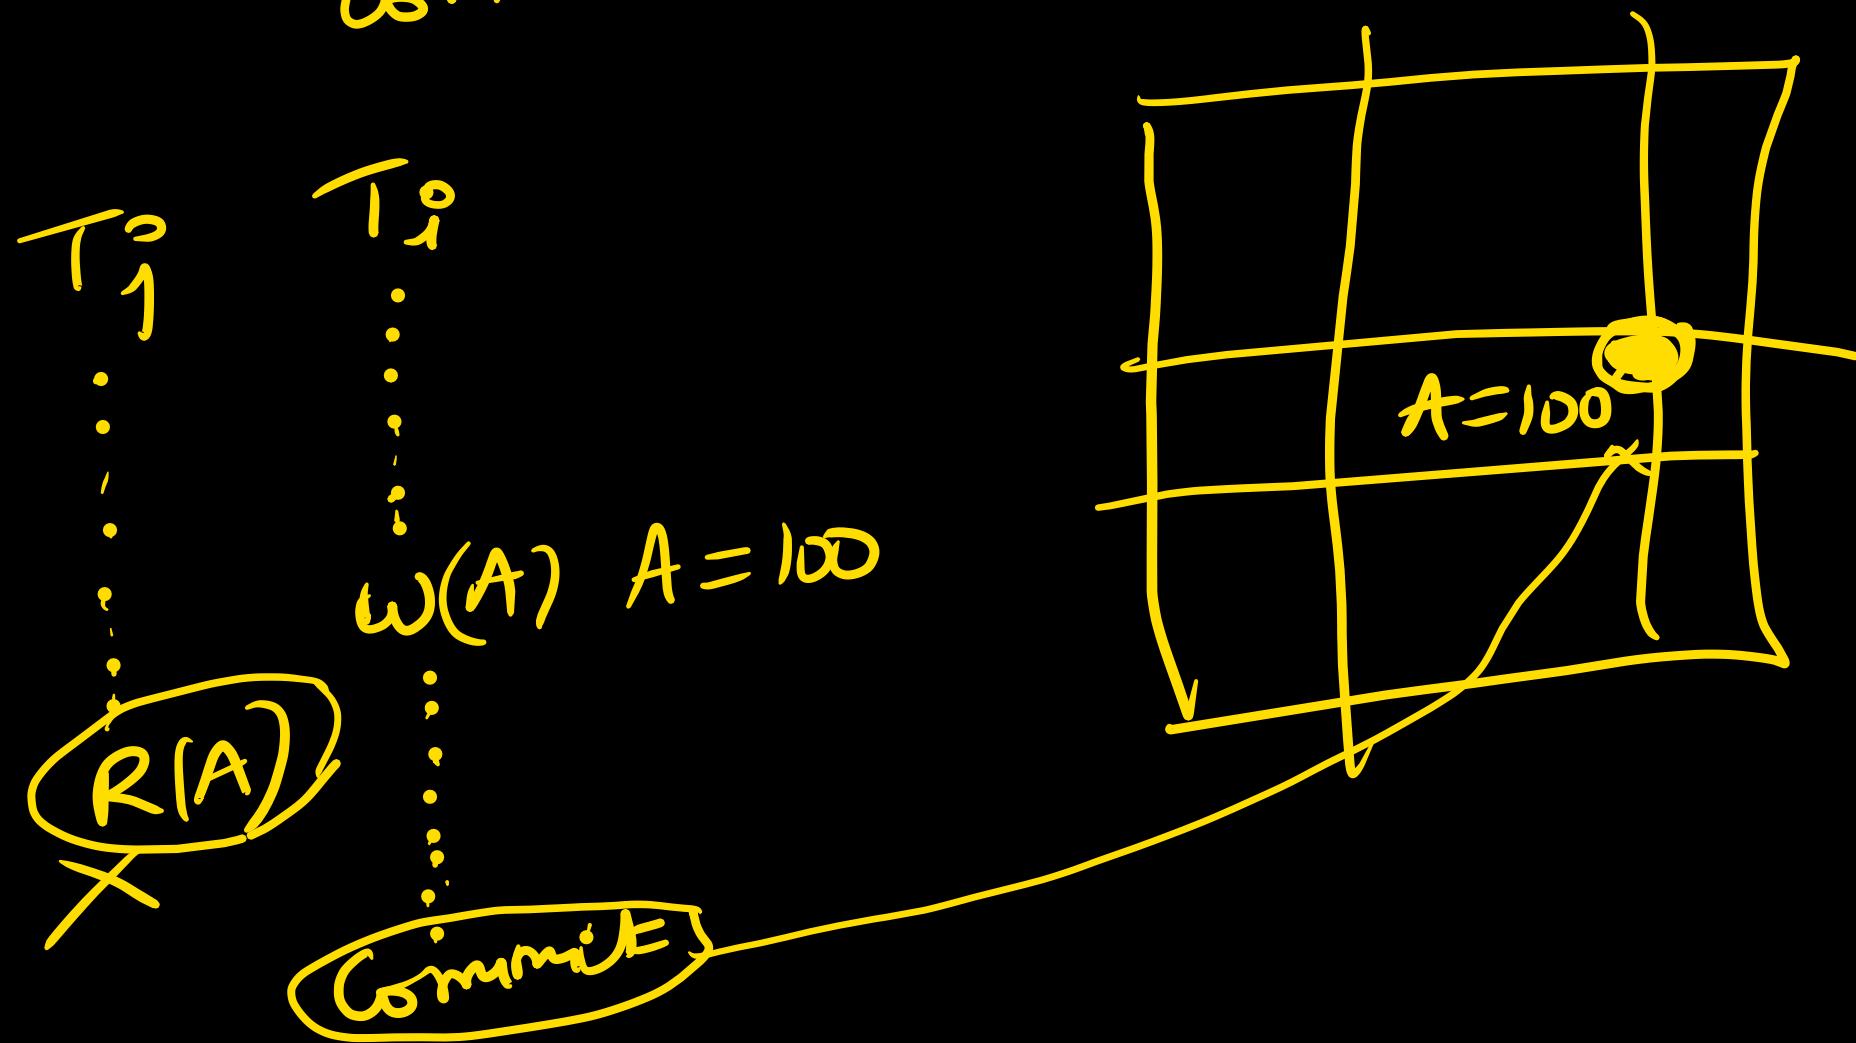
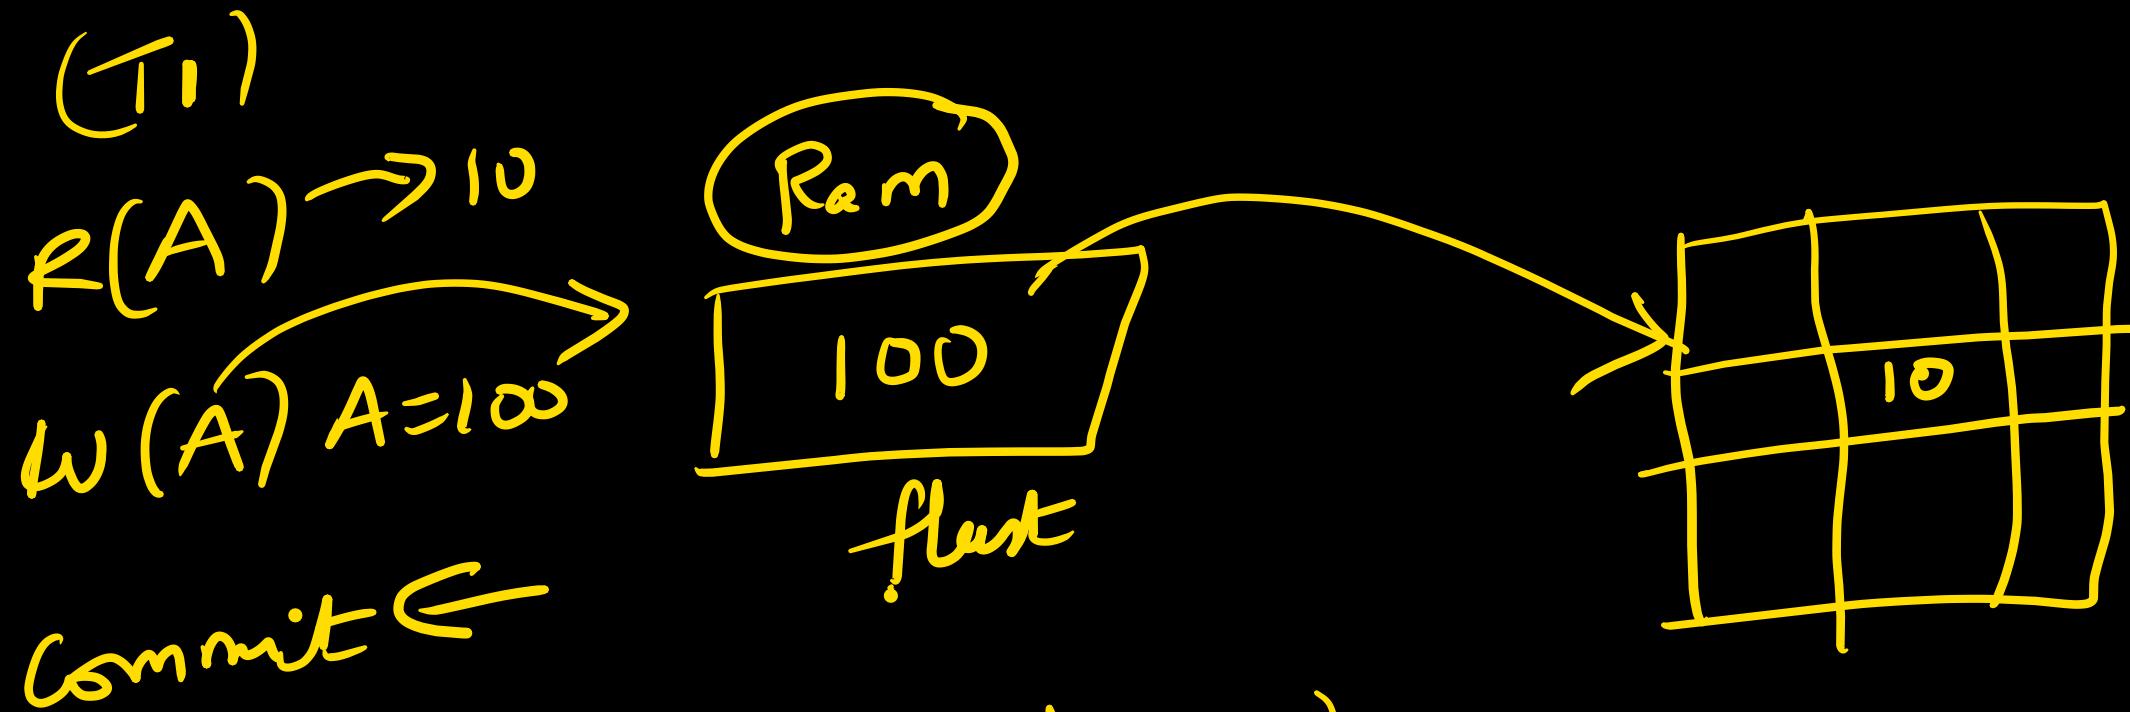
All modifications that will be done by because of commit. ✓
(Ex: removing of dirty flag ⚡ bit, flush flushing of data)

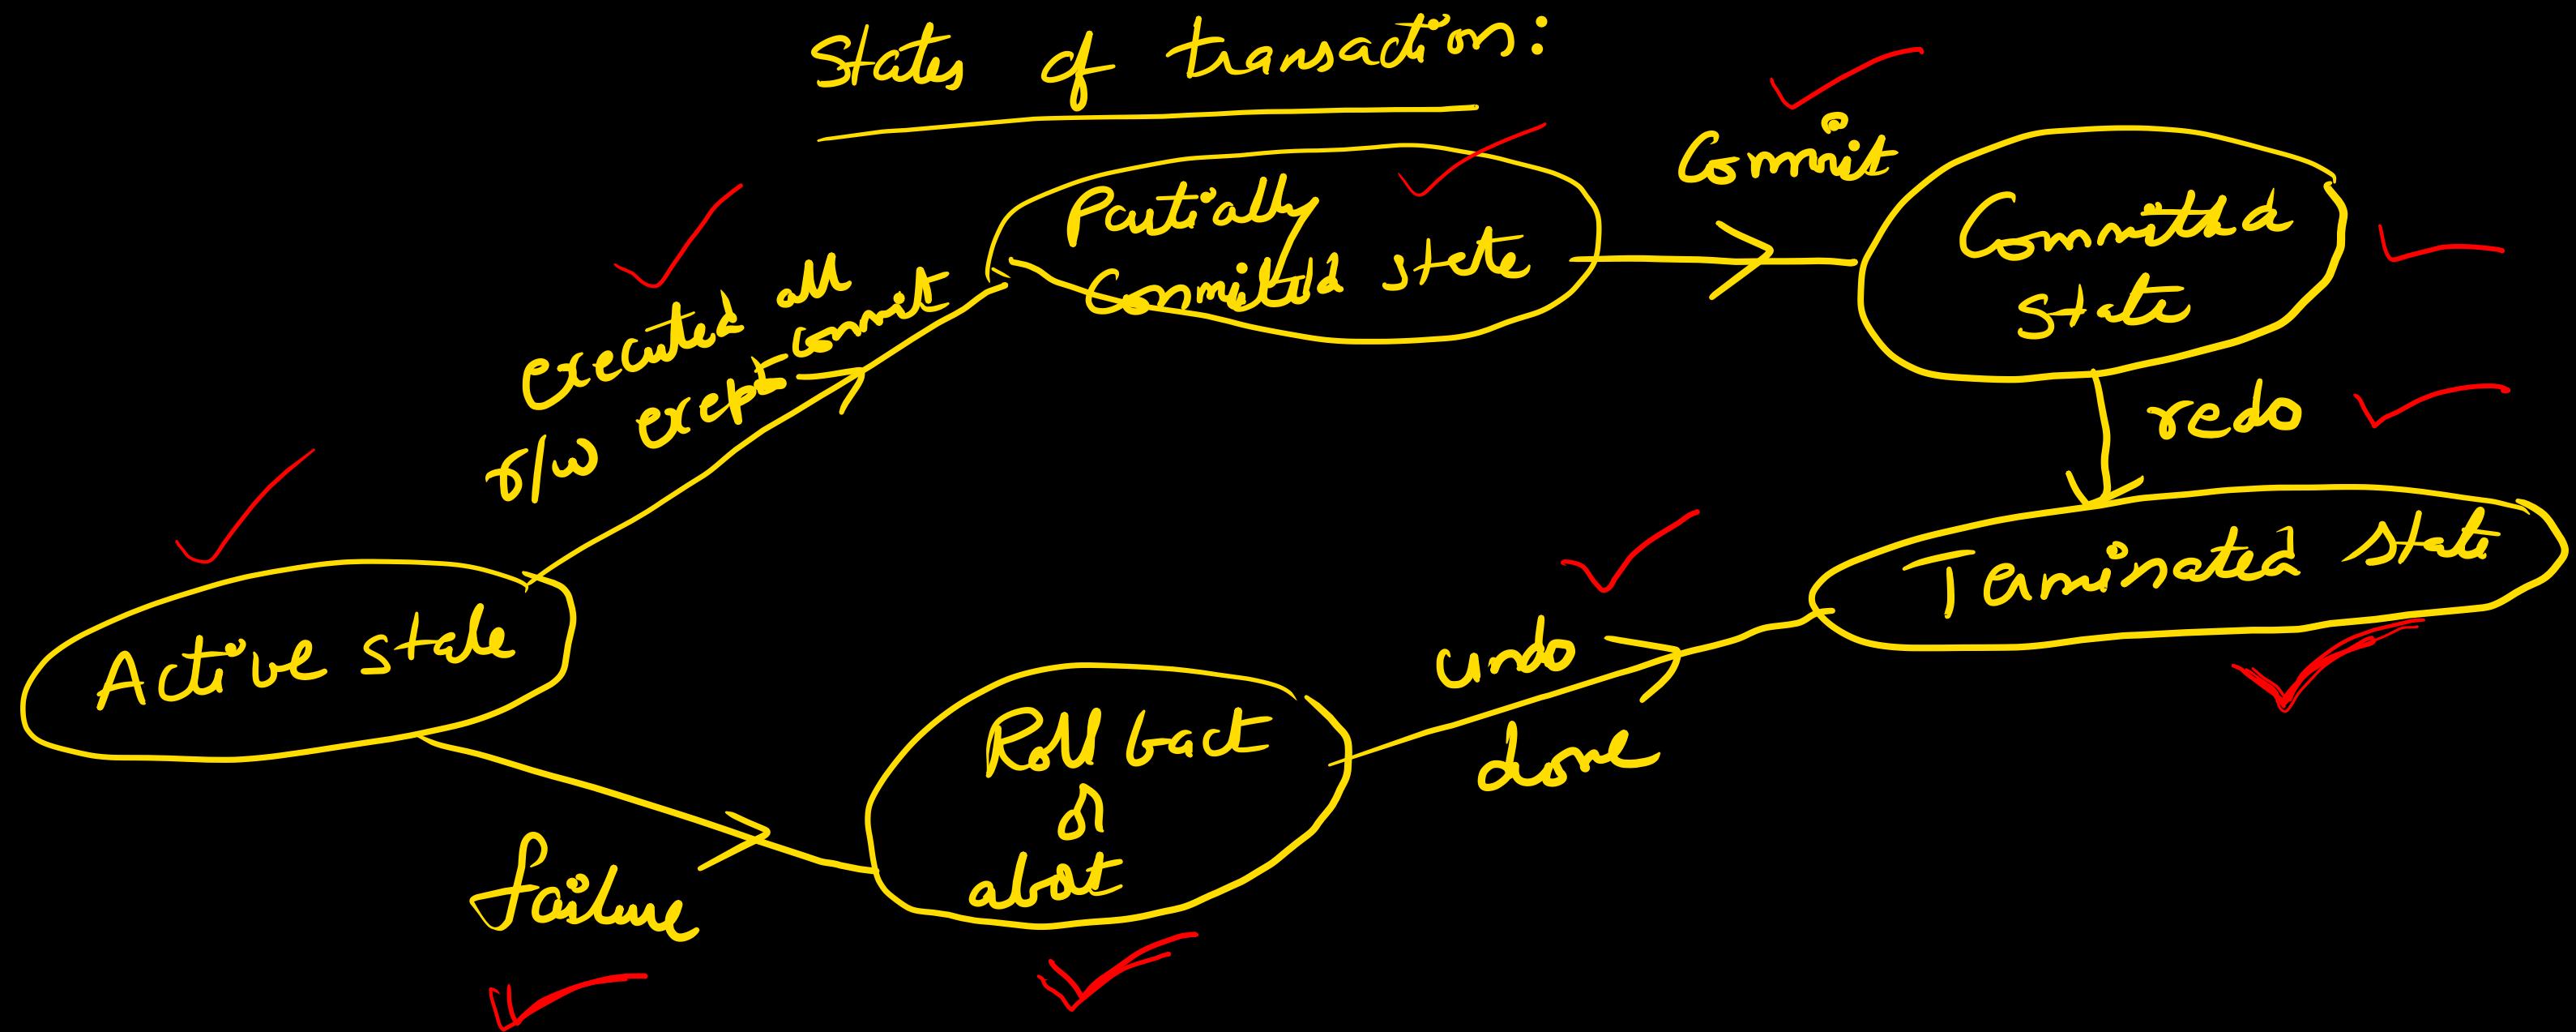
undo operation: ✓

All modification of OS file

because of transaction ~~lock~~





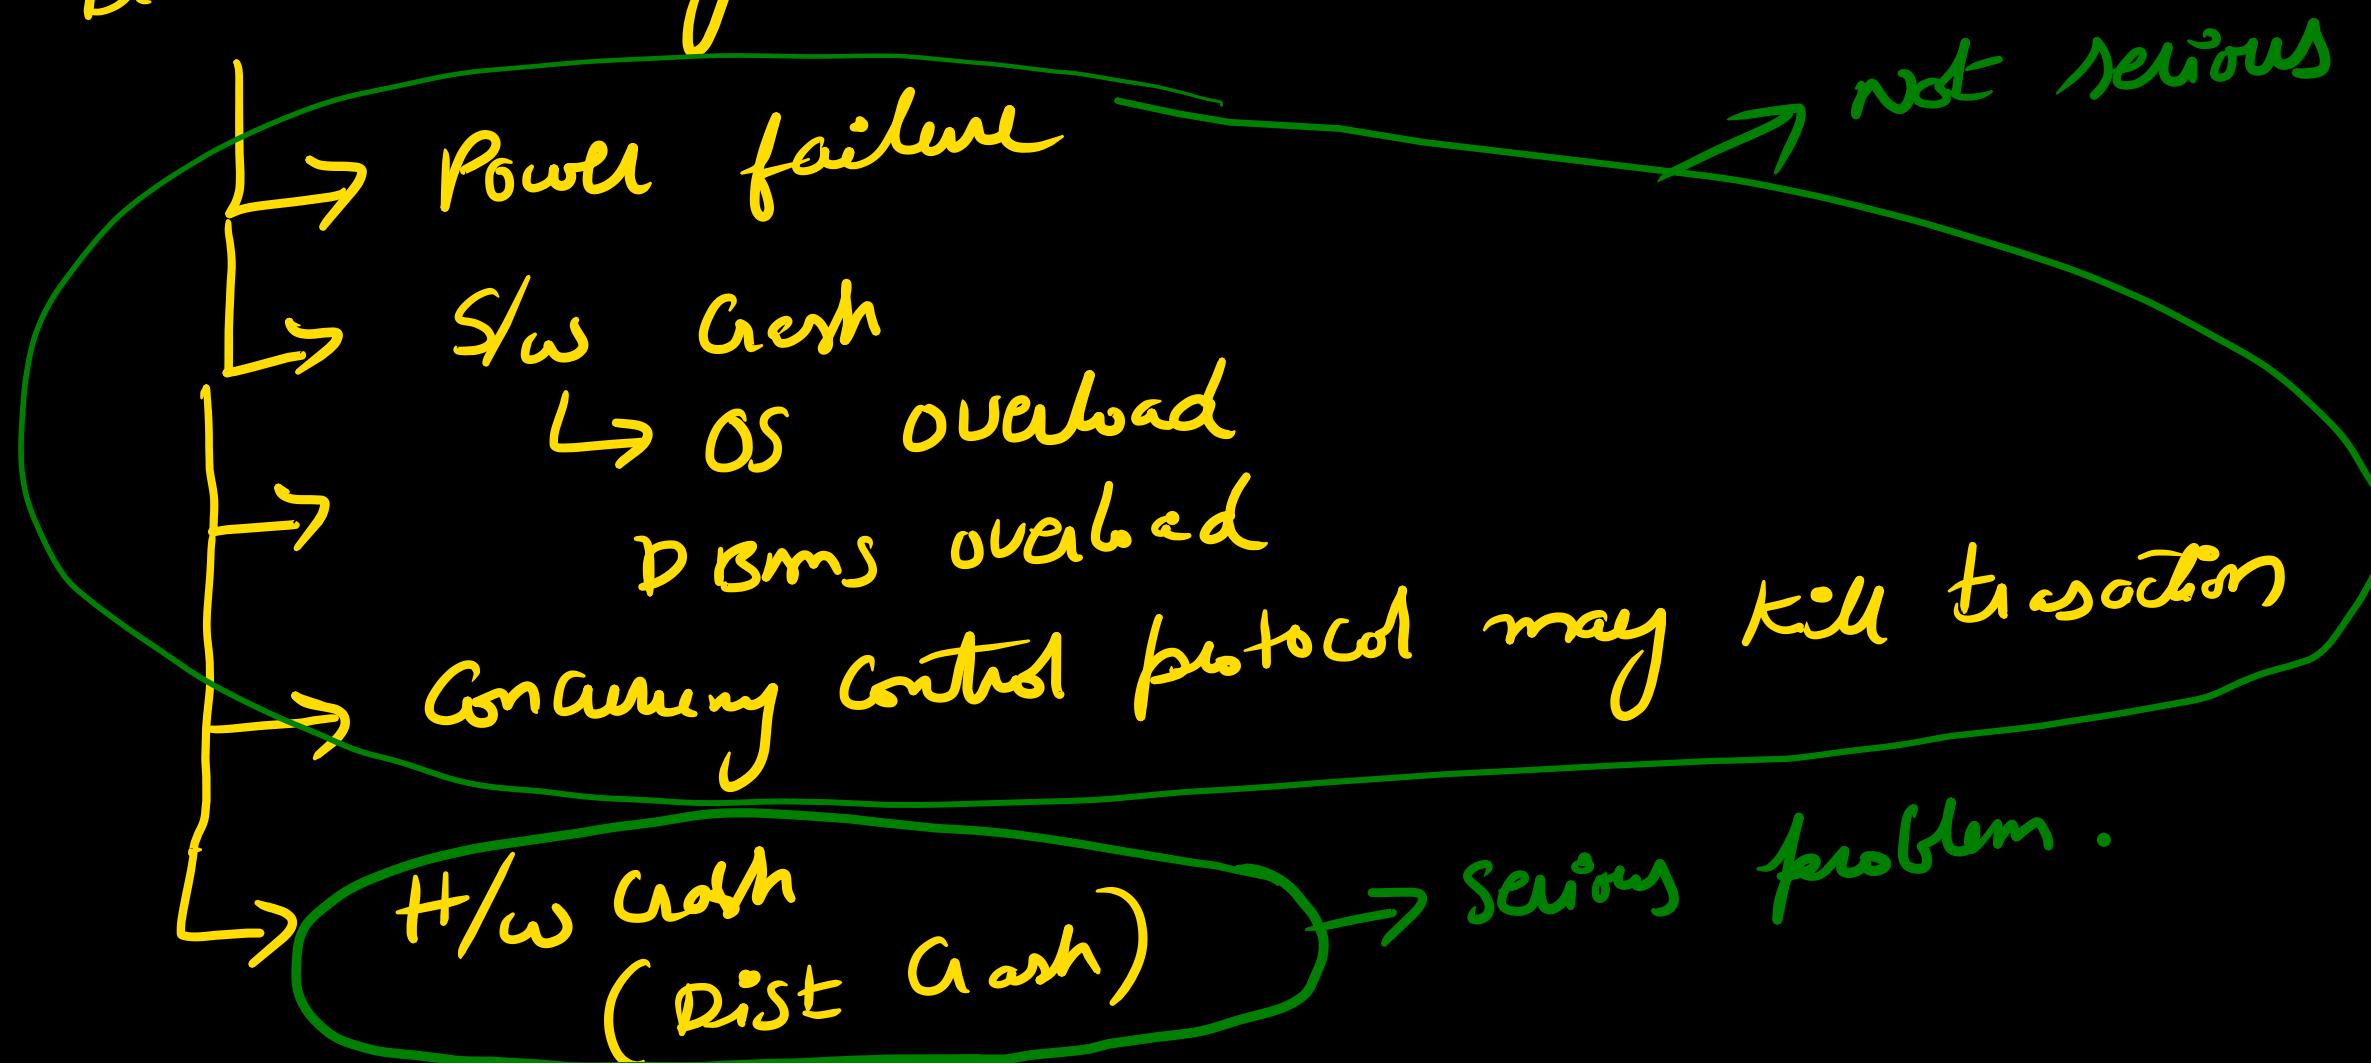


Atomicity ends here

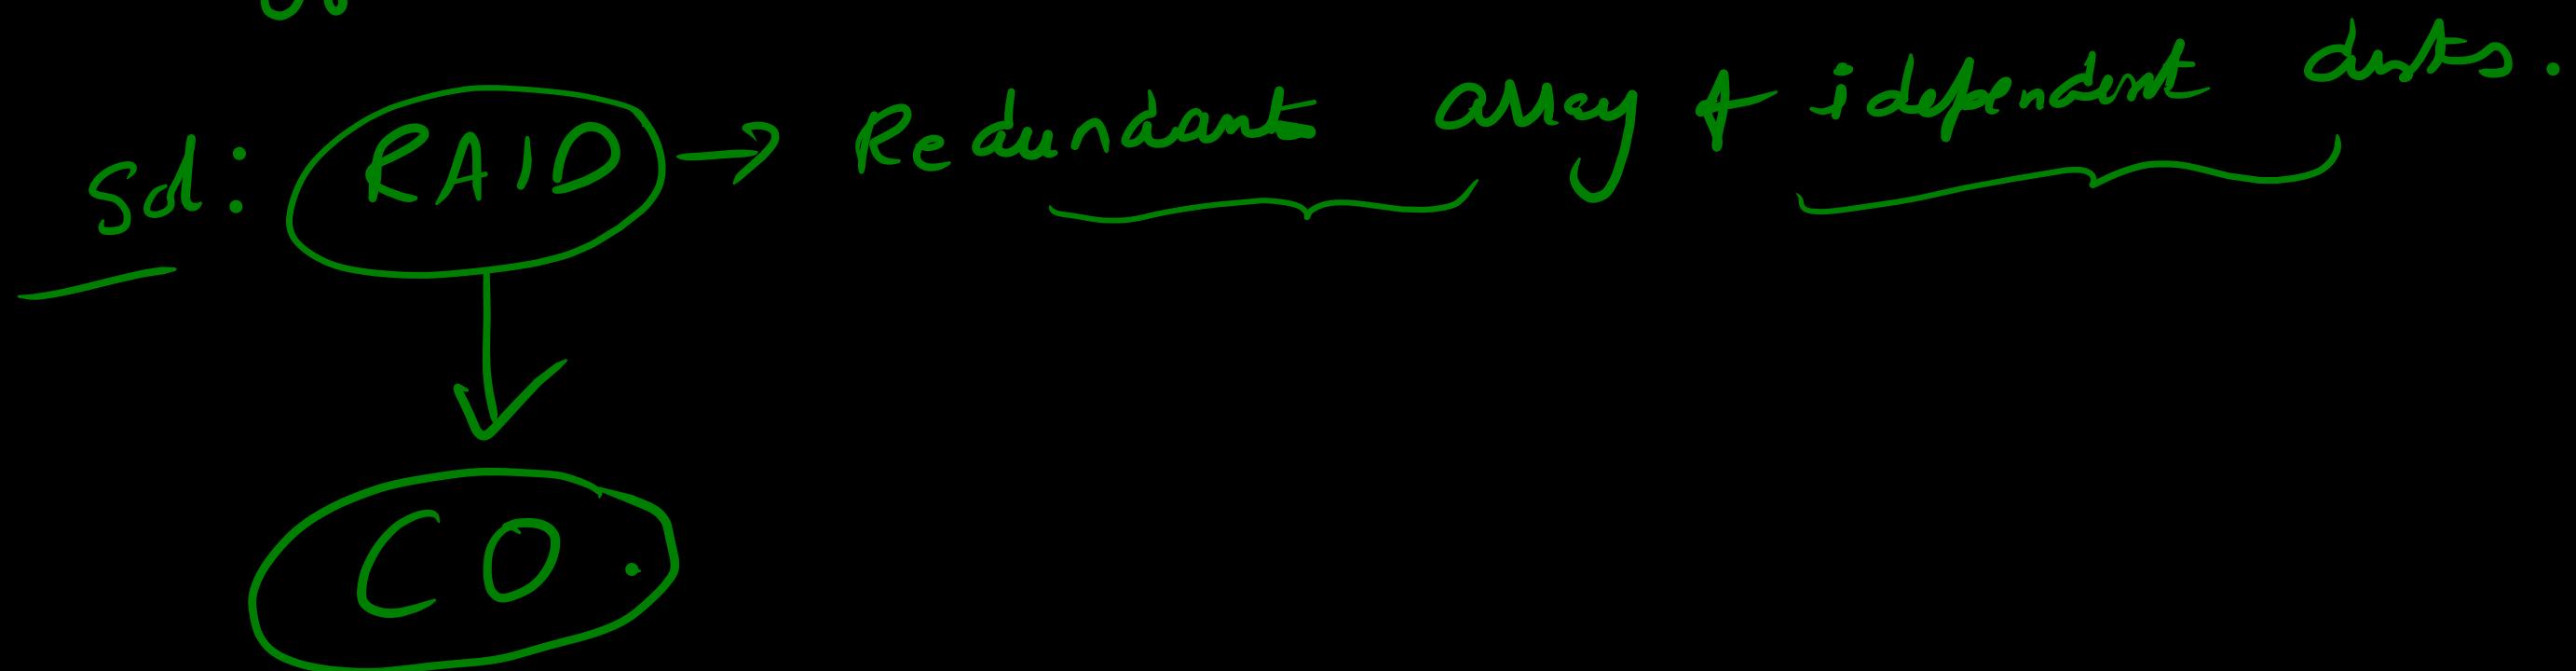
Durability :

Transaction should recover under any case of failure .

Transaction may fail because of following .



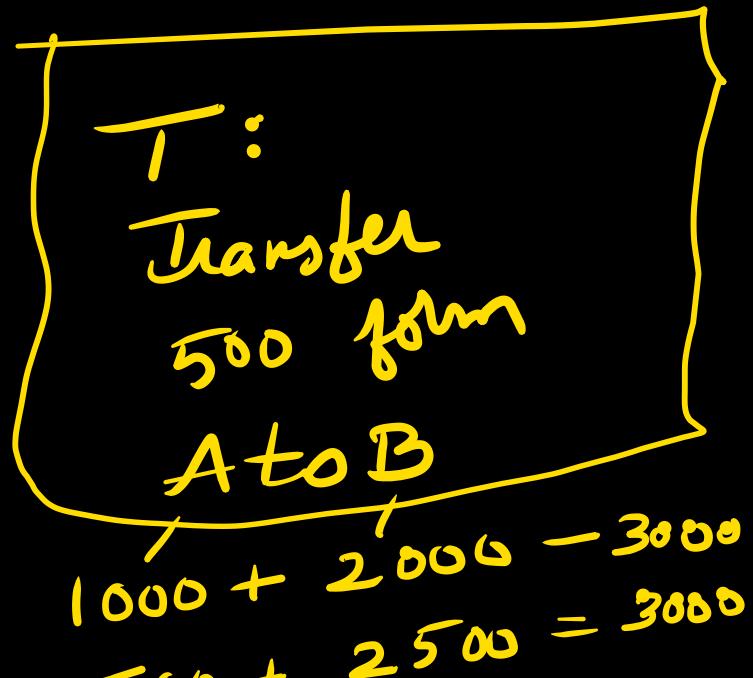
If disk fails → then transaction log will be lost.



Consistency:

This is user's responsibility.

User should check if the operations of a transaction is logically correct.



A+B before execution of Transaction



A+B after execution of transaction

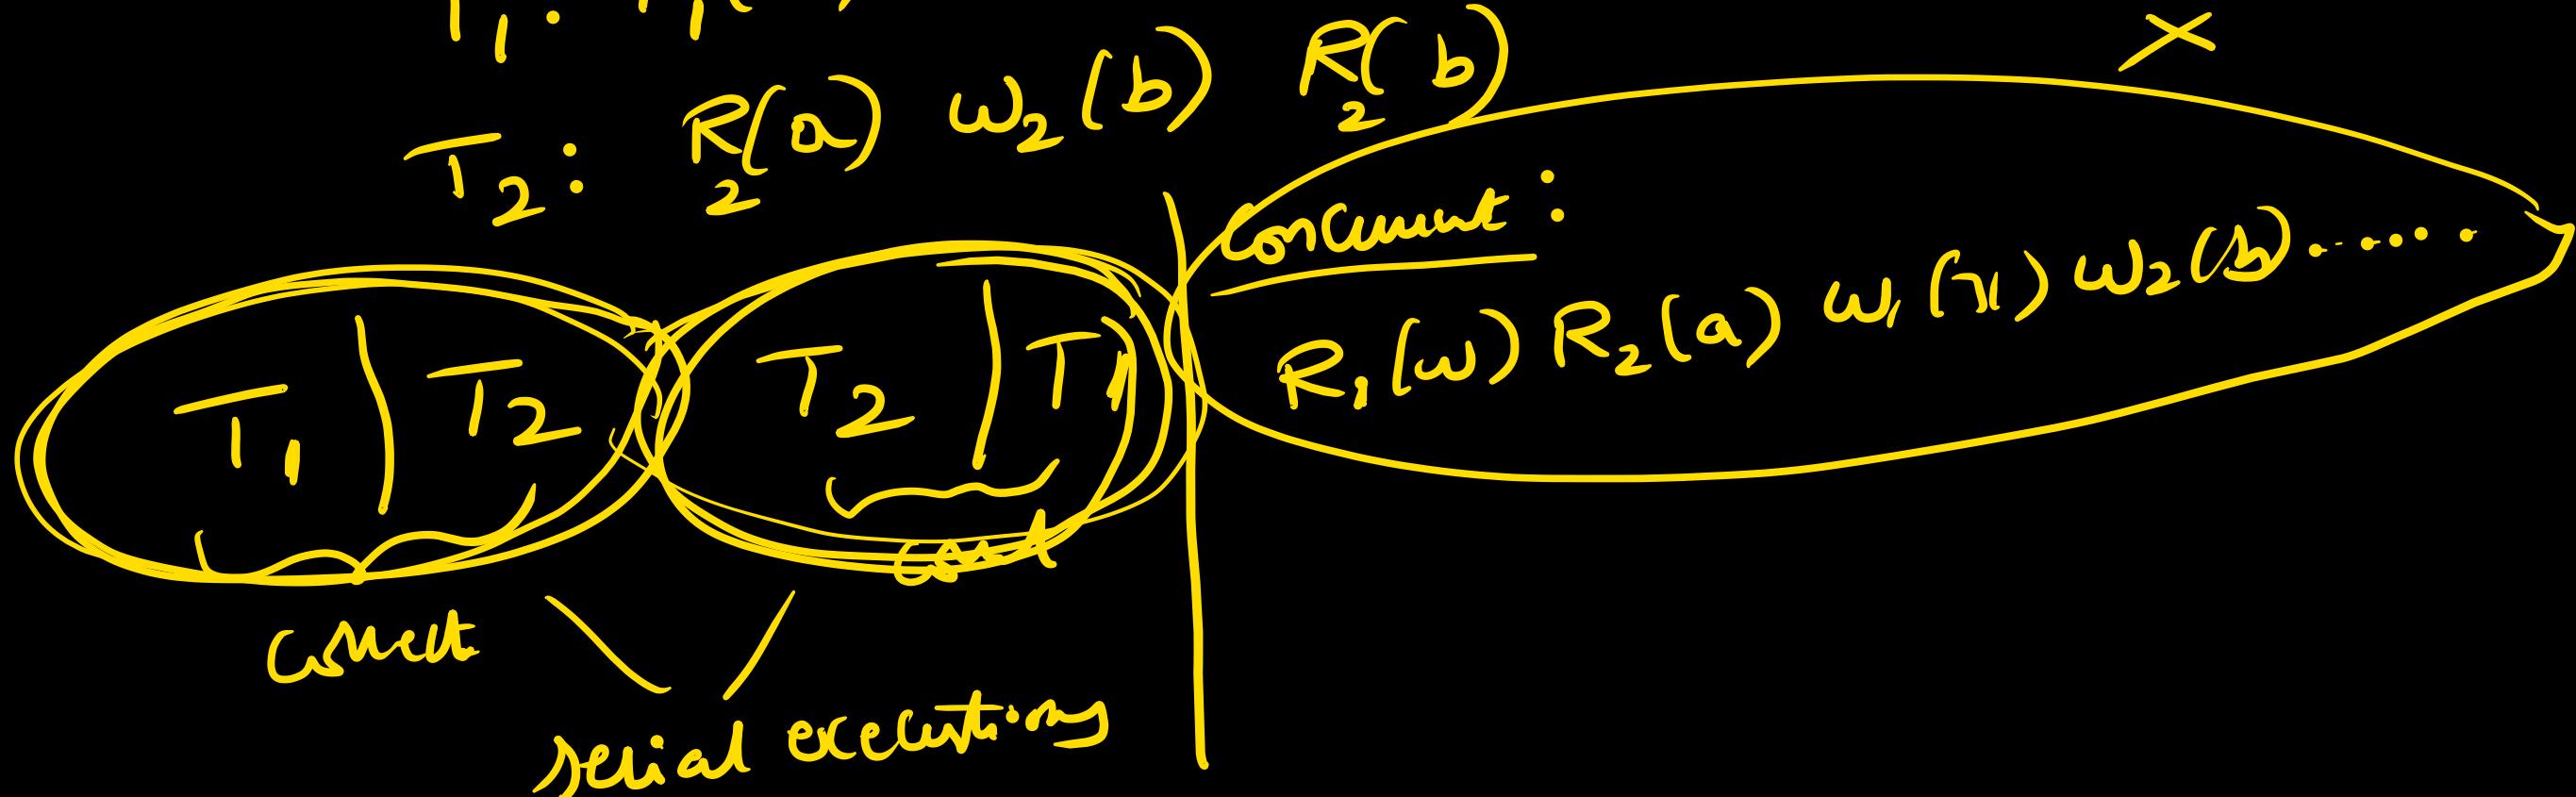
Consistency is over

Isolation: * * * * important fd gate

concurrent execution of two or more transactions must be equal
to result of any serial execution of transaction

$T_1: R_1(w) W_1(x) R_1(y) W_1(y)$

$T_2: R_2(a) W_2(b) R_2(b)$



Schedule: Time order creation sequence of two or more transactions.



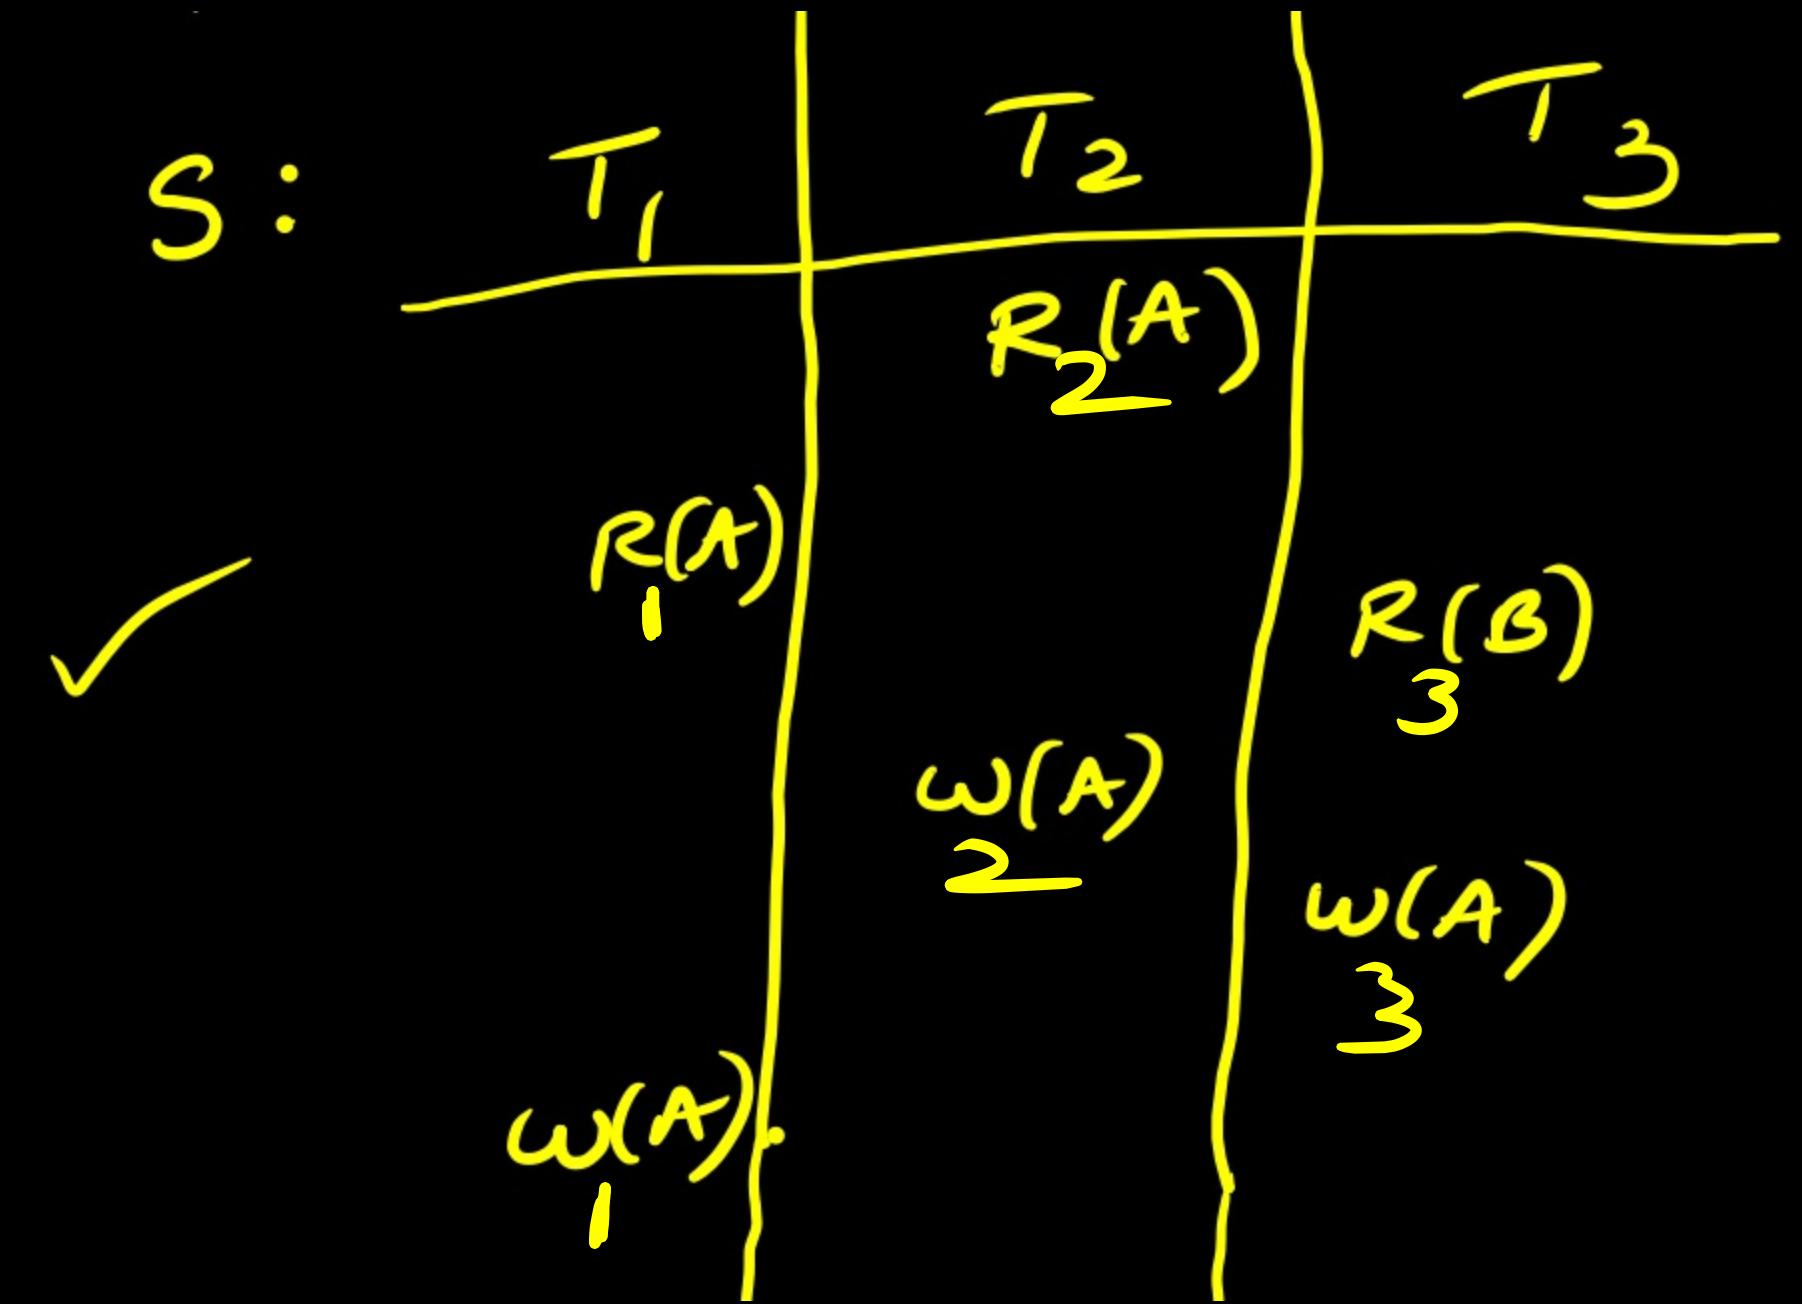
✓ S: $R_1(A), R_1(A), R_3(B), W_2(A), W_3(Q), W_1(A)$

(2)
2nd transaction

Trans: group of actions which are trying to achieve some task

Schedule is group of ~~interact~~ transactions

In a schedule,
there can be
any no of
transactions



Serial schedule:

Transactions should execute one after other without interleaving

T_1 : Transfer 500 from A to B

$[r_1(A) \quad w_1(A) \quad r_1(B) \quad w_1(B) \quad c_1]$

T_2 : Display A+B

$[r_2(A) \quad r_2(B) \quad c_2]$

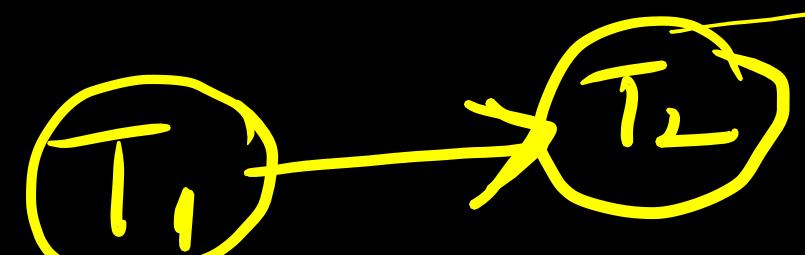
T_1 : Transfer 500 from A to B

$[r_1(A) \quad \omega_1(A) \quad r_1(B) \quad \omega_1(B) \quad c_1]$

T_2 : Display A + B

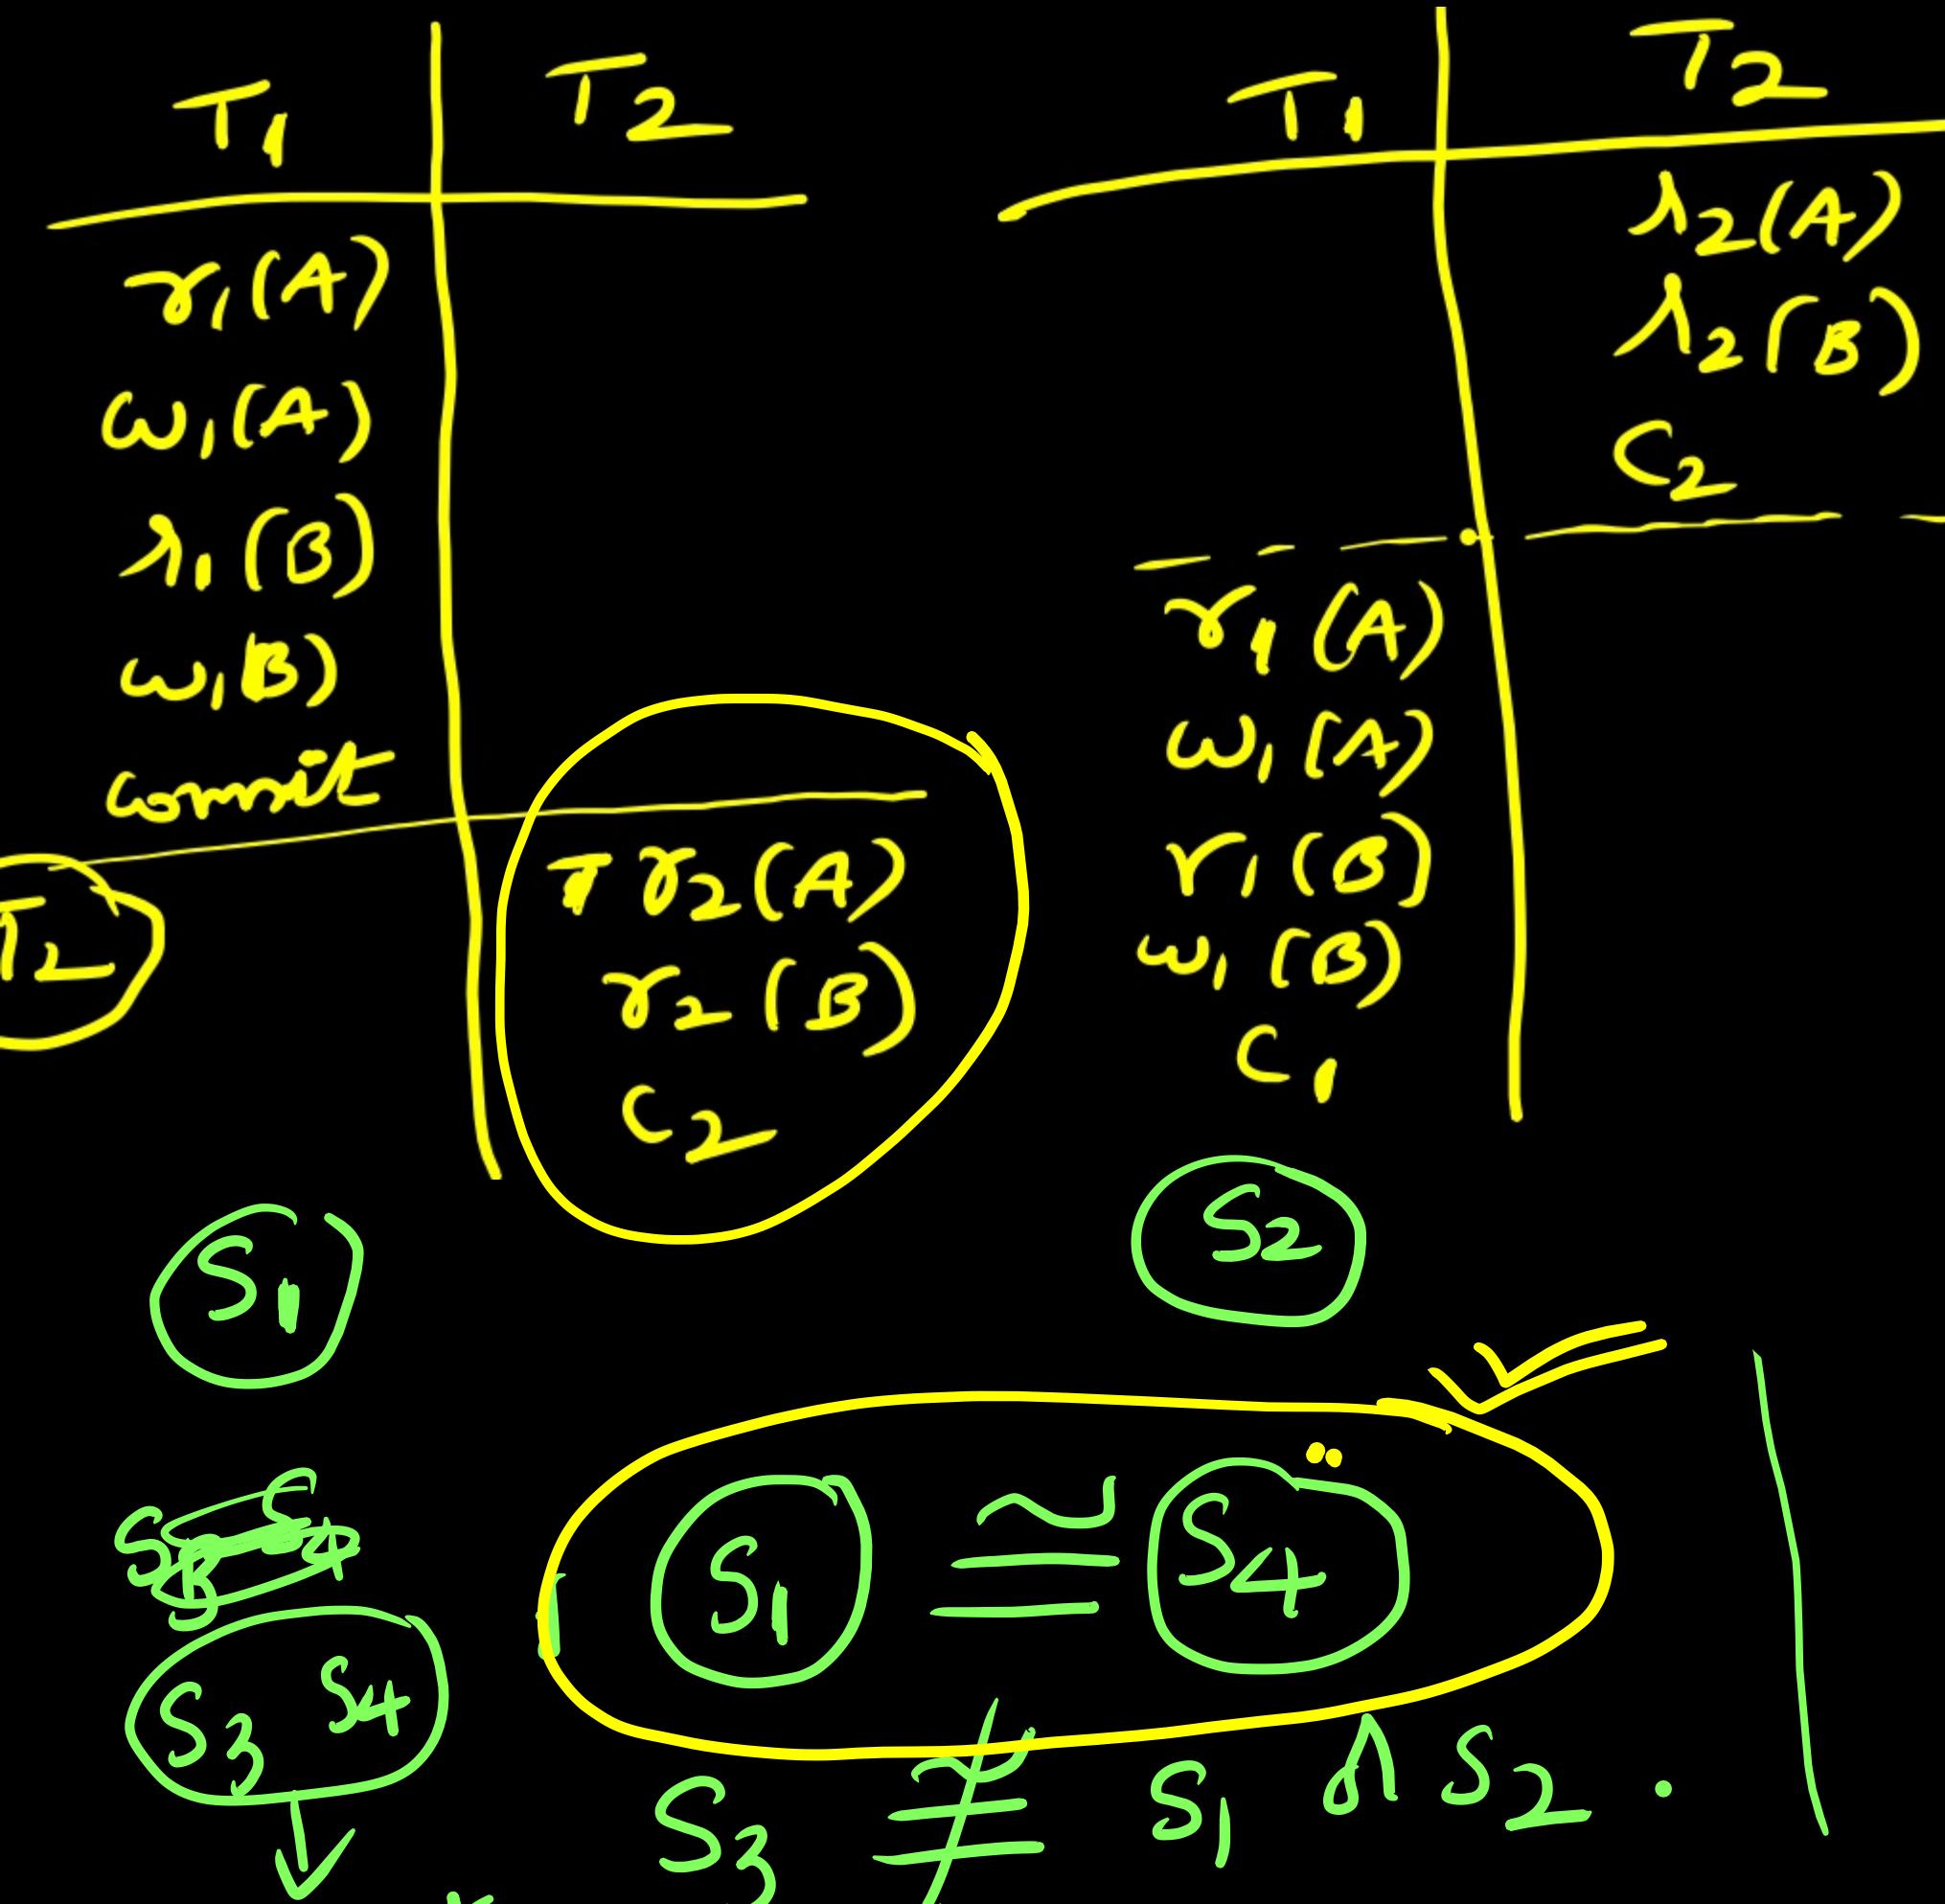
$[r_2(A) \quad r_2(B) \quad c_2]$

$\left. \begin{matrix} T_1, T_2 \\ T_2, T_1 \end{matrix} \right\}$ Serial schedule.

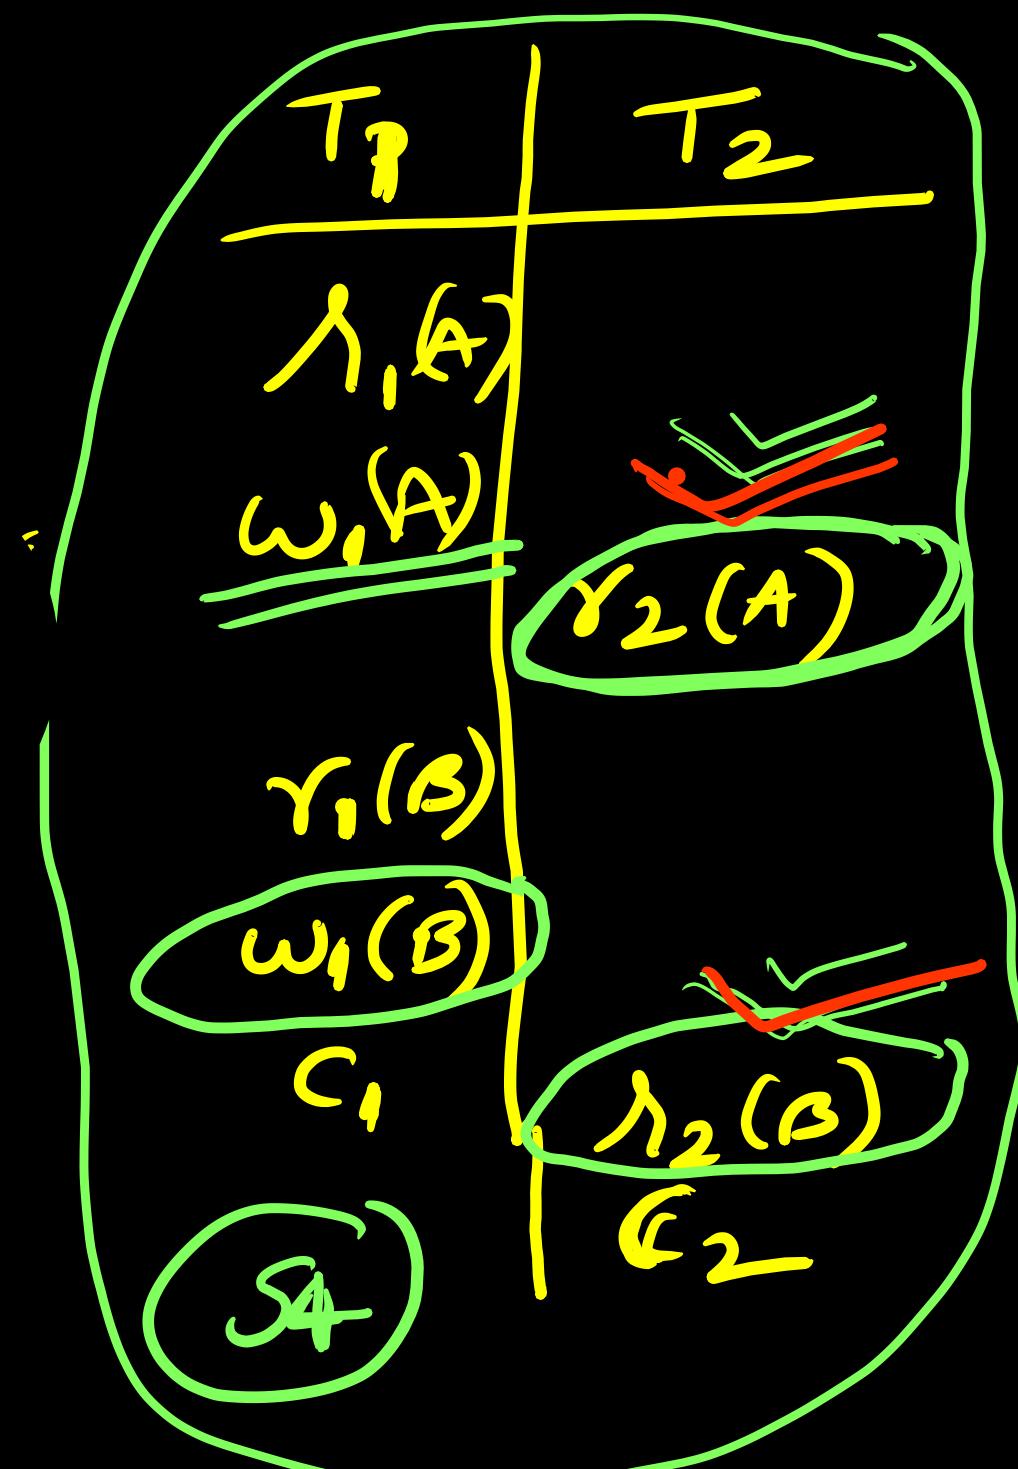


	T_1	T_2
	$r_1(A)$	
	$\omega_1(A)$	
	$r_1(B)$	
	$\omega_1(B)$	
commit		
		$r_2(A)$
		$\omega_2(A)$
		c_2

	T_1	T_2
		$\lambda_2(A)$
		$\lambda_2(B)$
		c_2
		-
	$\bar{r}_1(A)$	
	$\bar{\omega}_1(A)$	
	$\bar{r}_1(B)$	
	$\bar{\omega}_1(B)$	
	\bar{c}_1	



S_4 is equivalent
 to S_1
 $S_1 \cong S_4$
 $\therefore S_4$ is equivalent
 to a serial schedule
 S_1 .
 $\therefore S_4$ is called
 serializable



$T_1 \ T_2 \ \bar{T}_3$

$T_1 \ T_3 \ T_2$

$\bar{T}_2 \ T_1 \ T_3$

$T_2 \ T_3 \ T_1$

$\bar{T}_3 \ T_1 \ \bar{T}_2$

$\bar{T}_3 \ T_2 \ T_1$

$6(3!)$

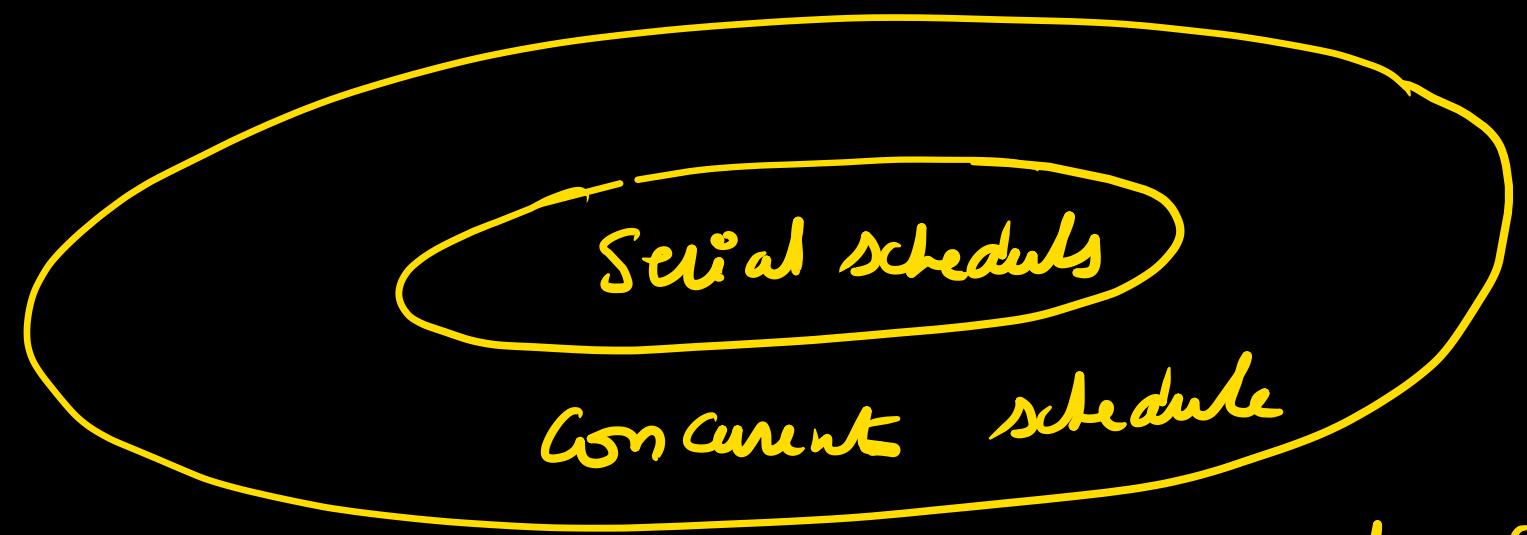
If there are 10 transactions, then
how many serial schedules will
be possible
 $(10!)$

adv: Data items in a serial schedule are accessed by one transaction at a time, there is no problem of inconsistency.

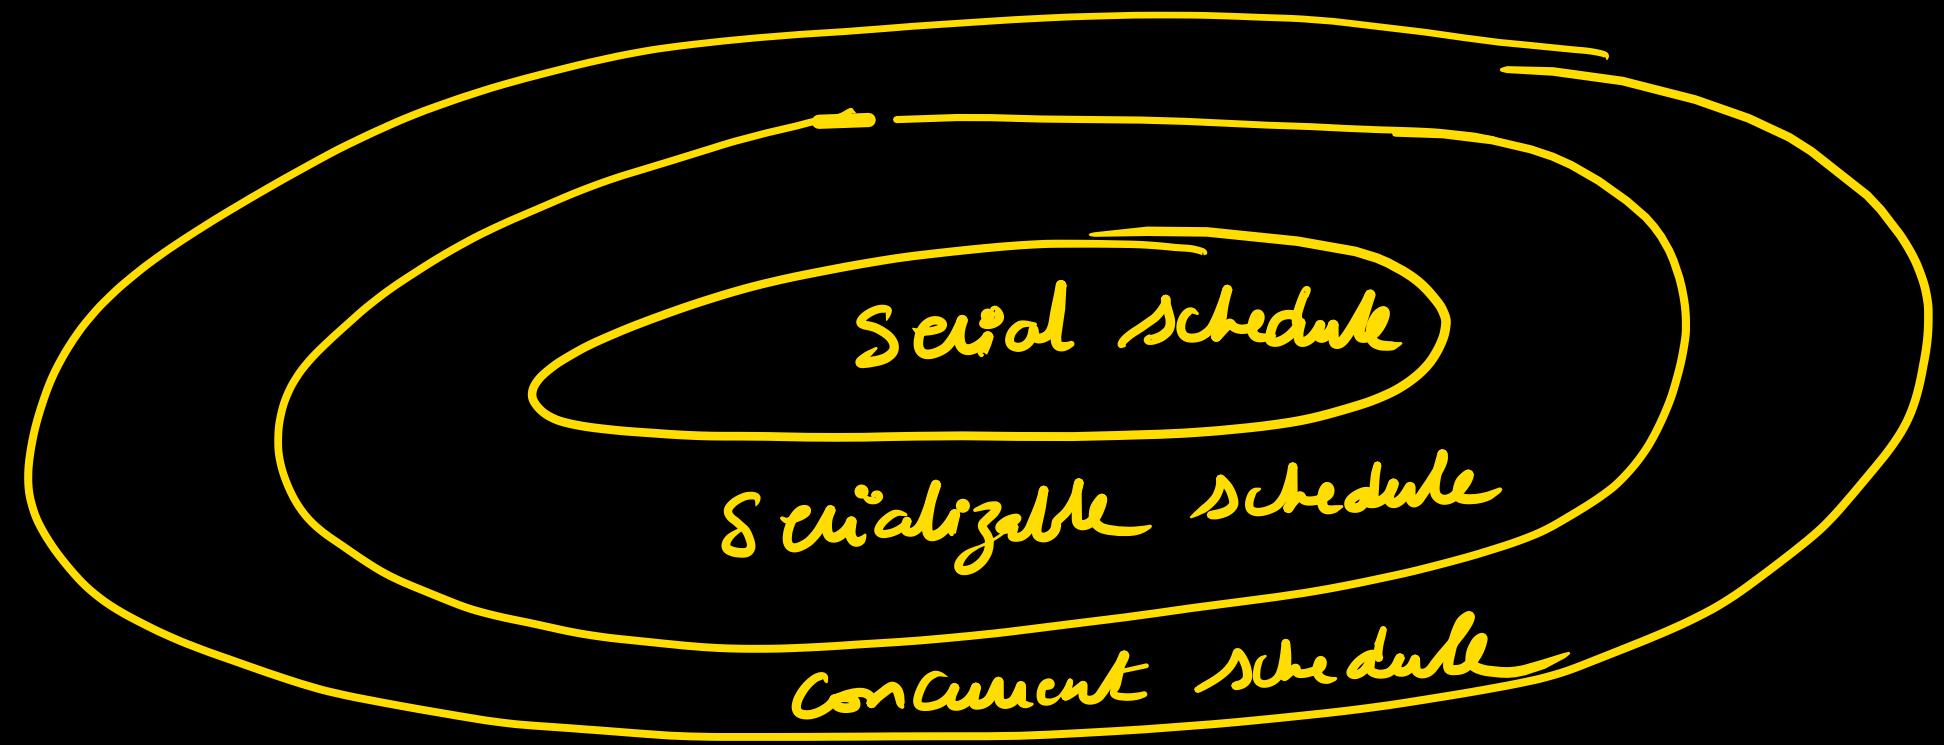
→ adv of serial schedules

dis: Co currency doesn't happen, less throughput, no sharing of resources.

Isolation: Concurrent schedule result must be equal to
result of any serial schedule.
These kind of schedules are also called serializable schedules.



Every serial schedule can also be called Concurrent Schedules
But all ~~serial schedules~~ there are schedules which
are not real serial.



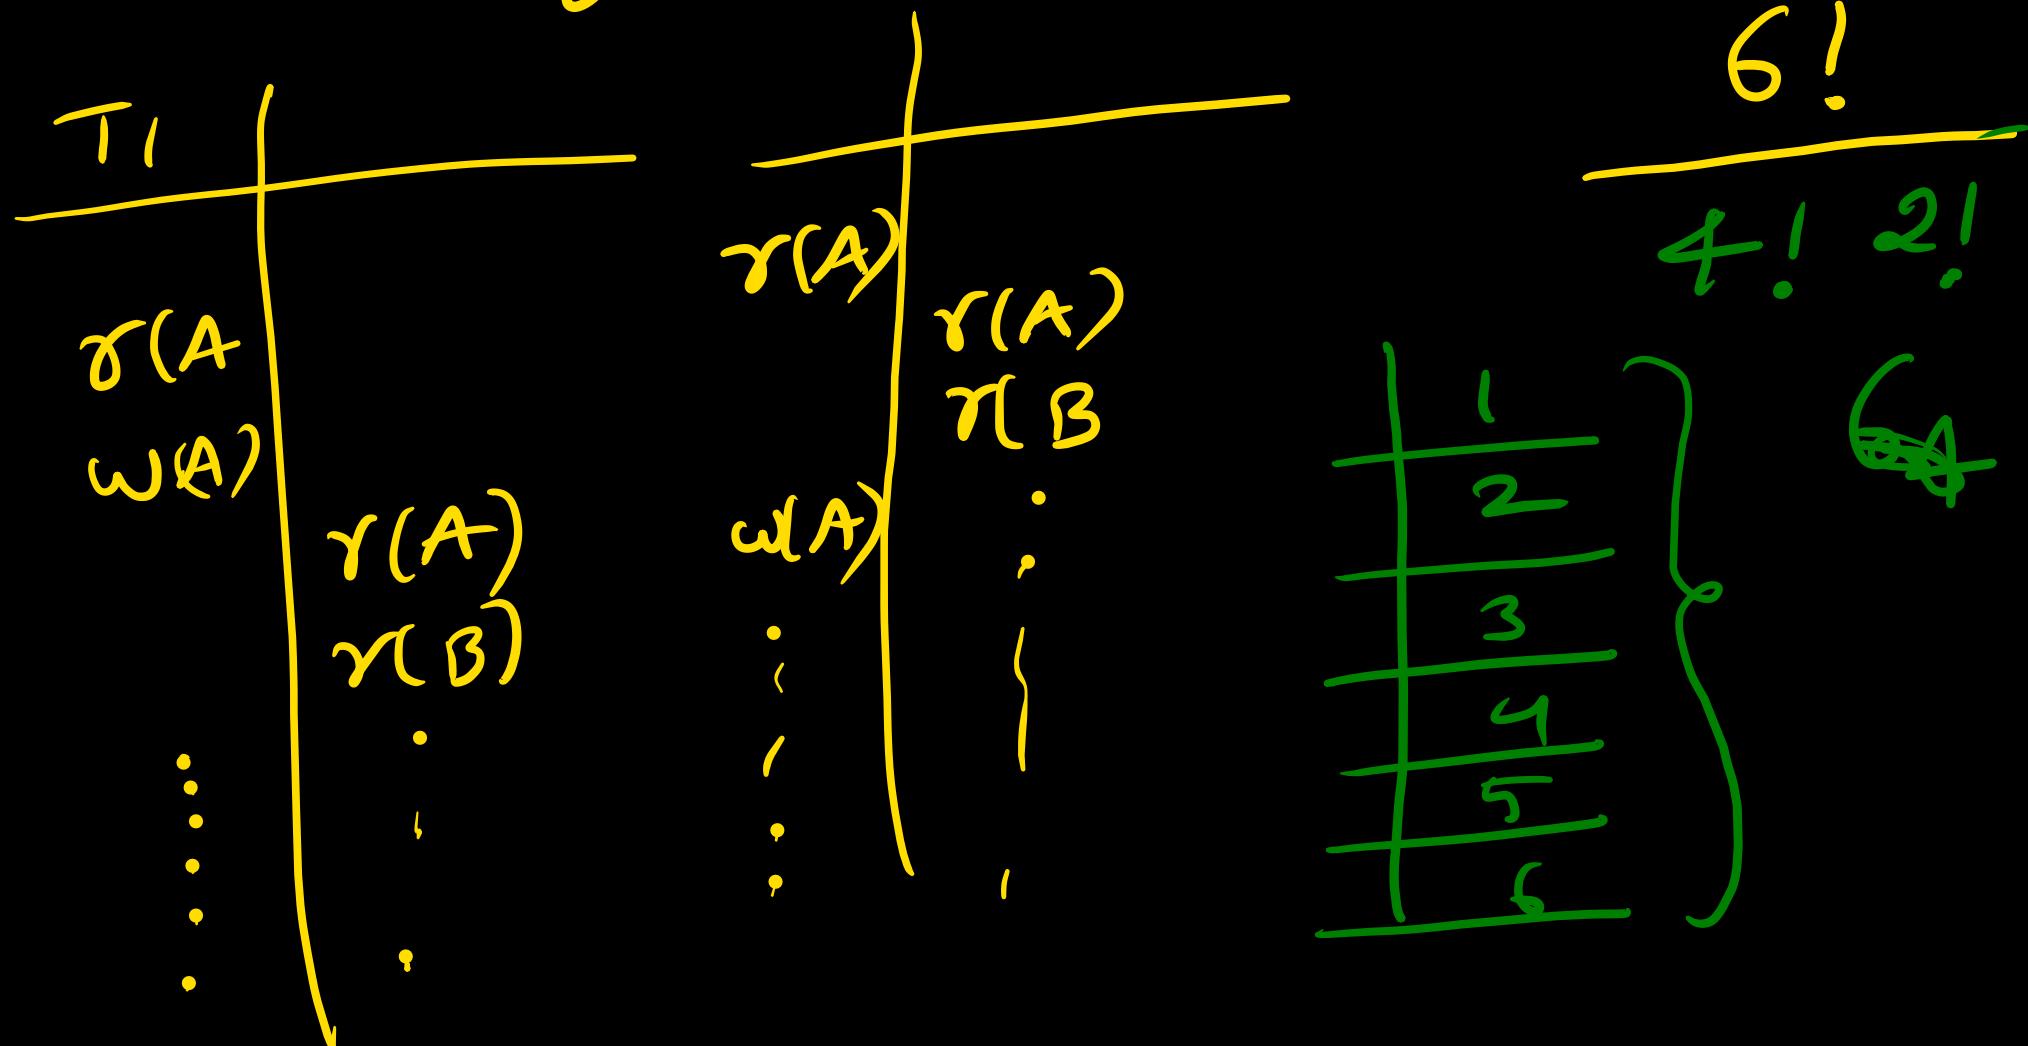
Concurrency control protocol:

It allows only the concurrent schedules which are equivalent to some serial schedule i.e. serializable schedules.

$T_1: \tau_1(A) \quad w_1(A) \quad r_1(B) \quad w_1(B)$ → no reassignment

$T_2: \tau_2(A) \quad r_2(B)$ → no reassignment

How many concurrent schedules are possible



$T_2: \tau_2(A) \quad r_2(B)$

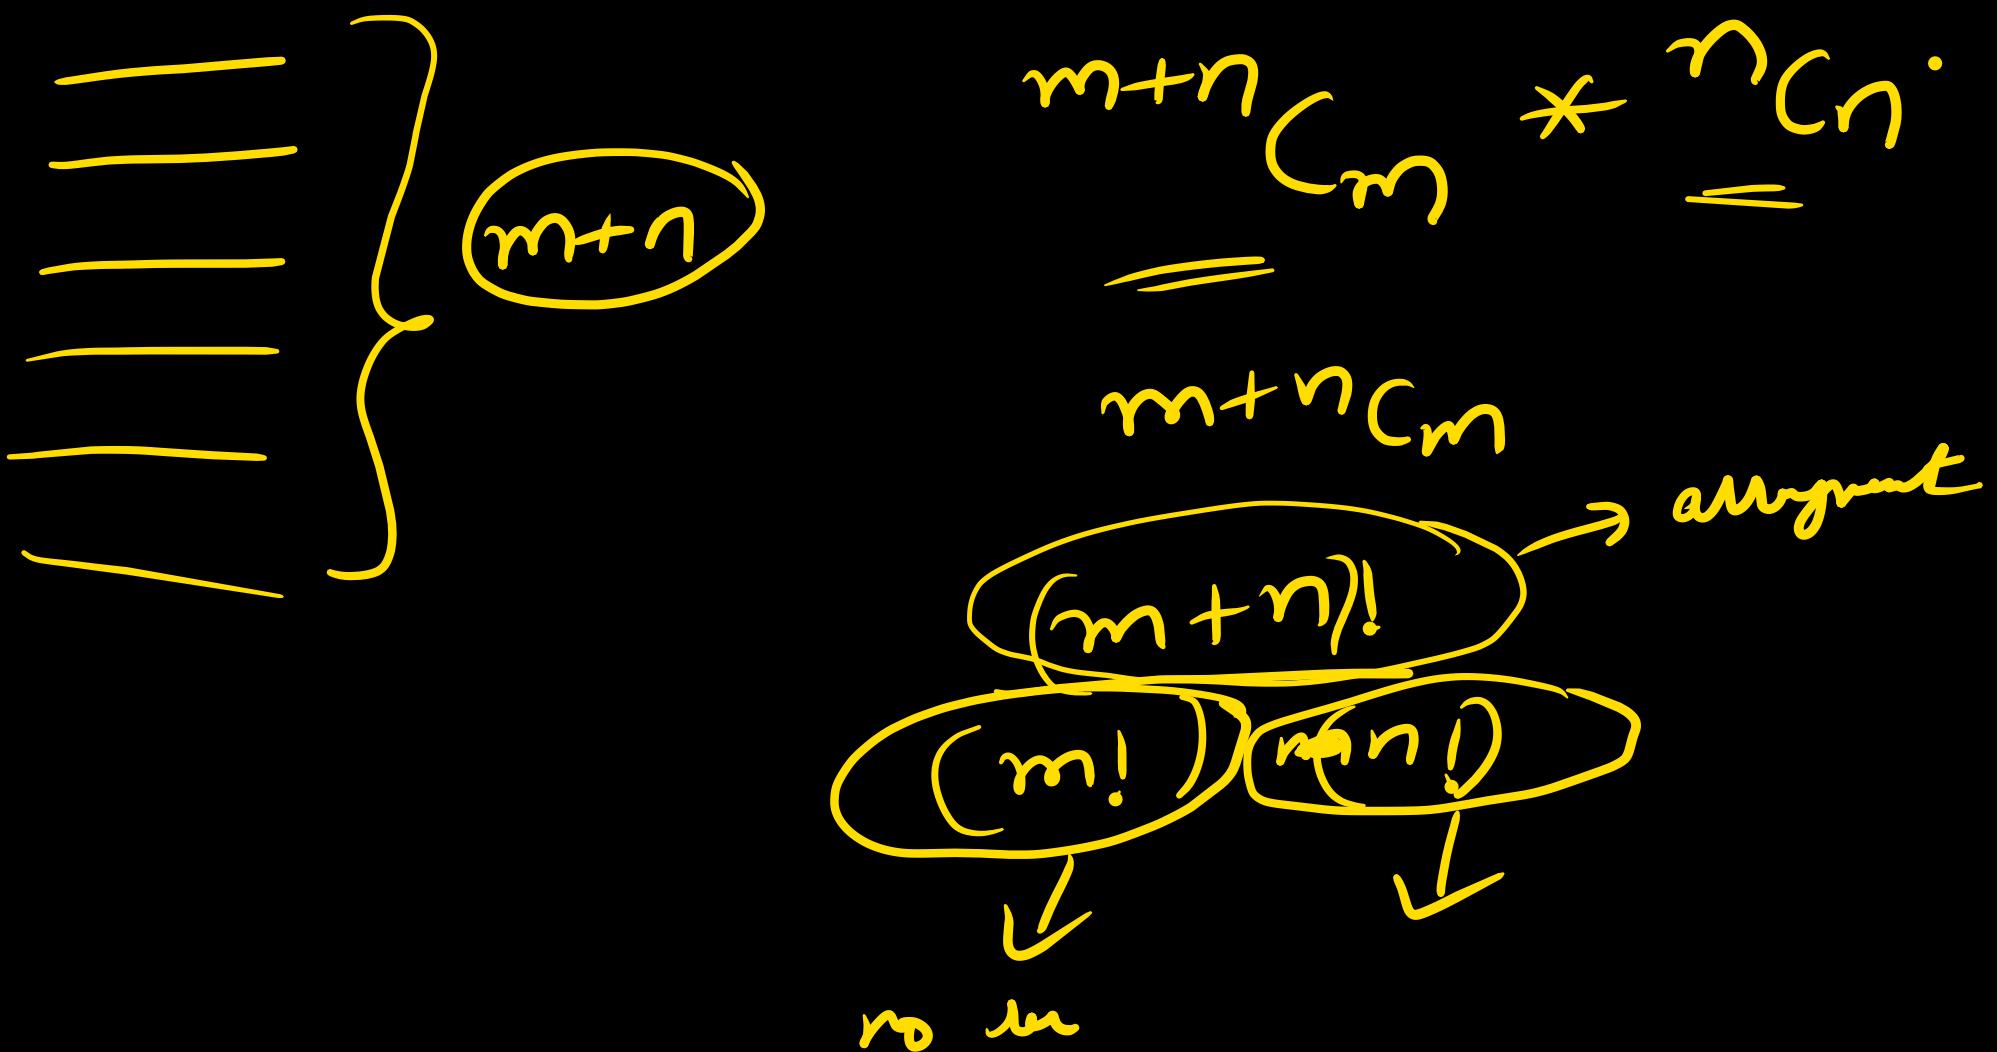
↓

$\tau_2(B) \quad r_2(A)$

$6C_4 \rightarrow T_1 \quad 2C_2 \rightarrow T_2$

$6C_4 \times 2C_2 = 6C_4.$

→ T, T_2 trans with $m & n$ operations each. How many concurrent schedules are possible?



T_1, T_2, T_3 are transactions with n, m, p operations each.
How many concurrent schedules are possible?

$$\frac{(n+m+p)!}{n! \ m! \ p!}$$

We have completed AAC properties.